

Project 2: Data Representations and Clustering

February 12, 2023

0.1 Module Loading

```
[75]: # modify this line to your local path
path = '/content/drive/MyDrive/Colab Notebooks/ECE219/Project2/'

[1]: import numpy as np
import pandas as pd
import random
import matplotlib.pyplot as plt
import matplotlib
%matplotlib inline
from IPython.core.pylabtools import figsize
import re
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
import string
from string import punctuation
from sklearn.decomposition import TruncatedSVD, NMF
from sklearn.utils.extmath import randomized_svd
from sklearn.metrics import auc, roc_curve, plot_roc_curve,
    ↳plot_confusion_matrix, accuracy_score, recall_score, precision_score,
    ↳f1_score
from sklearn import datasets, metrics, model_selection, svm
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.pipeline import Pipeline
from tempfile import mkdtemp
from joblib import Memory
from sklearn.multiclass import OneVsRestClassifier, OneVsOneClassifier
from sklearn.datasets import fetch_20newsgroups
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics.cluster import contingency_matrix, homogeneity_score,
    ↳completeness_score, v_measure_score, adjusted_rand_score,
    ↳adjusted_mutual_info_score
from sklearn.metrics import confusion_matrix
from scipy.optimize import linear_sum_assignment
```

```

!pip install umap-learn
import umap
!pip install hdbscan
import hdbscan
import pickle
import bz2
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.model_selection import train_test_split

```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting umap-learn

Downloading umap-learn-0.5.3.tar.gz (88 kB)

88.2/88.2 KB

9.0 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.8/dist-packages (from umap-learn) (1.21.6)

Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.8/dist-packages (from umap-learn) (1.0.2)

Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.8/dist-packages (from umap-learn) (1.7.3)

Requirement already satisfied: numba>=0.49 in /usr/local/lib/python3.8/dist-packages (from umap-learn) (0.56.4)

Collecting pynndescent>=0.5

Downloading pynndescent-0.5.8.tar.gz (1.1 MB)

1.1/1.1 MB

38.4 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from umap-learn) (4.64.1)

Requirement already satisfied: llvmlite<0.40,>=0.39.0dev0 in /usr/local/lib/python3.8/dist-packages (from numba>=0.49->umap-learn) (0.39.1)

Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-packages (from numba>=0.49->umap-learn) (57.4.0)

Requirement already satisfied: importlib-metadata in /usr/local/lib/python3.8/dist-packages (from numba>=0.49->umap-learn) (6.0.0)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.8/dist-packages (from pynndescent>=0.5->umap-learn) (1.2.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.22->umap-learn) (3.1.0)

Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.8/dist-packages (from importlib-metadata->numba>=0.49->umap-learn) (3.12.1)

Building wheels for collected packages: umap-learn, pynndescent

Building wheel for umap-learn (setup.py) ... done

```

Created wheel for umap-learn: filename=umap_learn-0.5.3-py3-none-any.whl
size=82829
sha256=d21bdbbbb4d585960b24595bfbf1f1caa6c48d54cd75cdcc6db5e4877d525a38
Stored in directory: /root/.cache/pip/wheels/a9/3a/67/06a8950e053725912e6a8c42
c4a3a241410f6487b8402542ea
Building wheel for pynndescent (setup.py) ... done
Created wheel for pynndescent: filename=pynndescent-0.5.8-py3-none-any.whl
size=55513
sha256=7e05b58c086f474870a25f5511f3a656c1146c71fad07984fc699e23077109bf
Stored in directory: /root/.cache/pip/wheels/1c/63/3a/29954bca1a27ba100ed8c279
73a78cb71b43dc67aed62e80c3
Successfully built umap-learn pynndescent
Installing collected packages: pynndescent, umap-learn
Successfully installed pynndescent-0.5.8 umap-learn-0.5.3
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Collecting hdbscan
  Downloading hdbscan-0.8.29.tar.gz (5.2 MB)
      5.2/5.2 MB
87.1 MB/s eta 0:00:00
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Requirement already satisfied: scikit-learn>=0.20 in
/usr/local/lib/python3.8/dist-packages (from hdbscan) (1.0.2)
Requirement already satisfied: joblib>=1.0 in /usr/local/lib/python3.8/dist-
packages (from hdbscan) (1.2.0)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.8/dist-
packages (from hdbscan) (1.21.6)
Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.8/dist-
packages (from hdbscan) (1.7.3)
Requirement already satisfied: cython>=0.27 in /usr/local/lib/python3.8/dist-
packages (from hdbscan) (0.29.33)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.20->hdbscan)
(3.1.0)
Building wheels for collected packages: hdbscan
  Building wheel for hdbscan (pyproject.toml) ... done
  Created wheel for hdbscan: filename=hdbscan-0.8.29-cp38-cp38-linux_x86_64.whl
size=3773945
sha256=c6ba2a629c1df8a860bd5b00c970d289520fc1838accc963271ea407a4c79e4e
Stored in directory: /root/.cache/pip/wheels/76/06/48/527e038689c581cc9e519c73
840efdc7473805149e55bd7ffd
Successfully built hdbscan
Installing collected packages: hdbscan
Successfully installed hdbscan-0.8.29

```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
import sys
sys.path.append('/content/drive/MyDrive/Colab Notebooks/ECE219/Project2')
from plotmat import plot_mat
```

Mounted at /content/drive

```
[3]: # helper code
!pip install torch
!pip install torchvision
import torch
import torch.nn as nn
from torchvision import transforms, datasets
from torch.utils.data import DataLoader, TensorDataset
from tqdm import tqdm
import requests
import os
import tarfile
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.base import TransformerMixin
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: torch in /usr/local/lib/python3.8/dist-packages (1.13.1+cu116)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torch) (4.4.0)
Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: torchvision in /usr/local/lib/python3.8/dist-packages (0.14.1+cu116)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from torchvision) (2.25.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from torchvision) (1.21.6)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torchvision) (4.4.0)
Requirement already satisfied: torch==1.13.1 in /usr/local/lib/python3.8/dist-packages (from torchvision) (1.13.1+cu116)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.8/dist-packages (from torchvision) (7.1.2)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (1.24.3)
Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (4.0.0)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (2.10)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (2022.12.7)

0.2 Question 1

Report the dimensions of the TF-IDF matrix you obtain.

Ans:

the dimensions of the TF-IDF matrix is (7882, 23522).

```
[ ]: categories = ['comp.graphics', 'comp.os.ms-windows.misc', 'comp.sys.ibm.pc.
↳ hardware',
                  'comp.sys.mac.hardware', 'rec.autos', 'rec.motorcycles', 'rec.sport.
↳ baseball', 'rec.sport.hockey']
dataset = fetch_20newsgroups(subset = 'all', categories = categories, shuffle =
↳ True, random_state = 0, remove=('headers', 'footers'))

vec = CountVectorizer(stop_words='english', min_df=3)
tfidf = TfidfTransformer()

X_vec = vec.fit_transform(dataset.data)
X_tfidf = tfidf.fit_transform(X_vec)
print("TFIDF dataset shape:", X_tfidf.shape)
```

TFIDF dataset shape: (7882, 23522)

0.3 Question 2

Report the contingency table of your clustering result. You may use the provided plotmat.py to visualize the matrix. Does the contingency matrix have to be square-shaped?

Ans:

The contingency table is

| | |
|------|------|
| 3232 | 671 |
| 54 | 3925 |

The contingency matrix doesn't have to be square-shaped but it's usually square-shaped. It has normally same rows and columns because the reference data and the map should have same dimensions (categories). However, it's not necessary, a category can exist in the reference data but doesn't exist in the map, and vice versa.

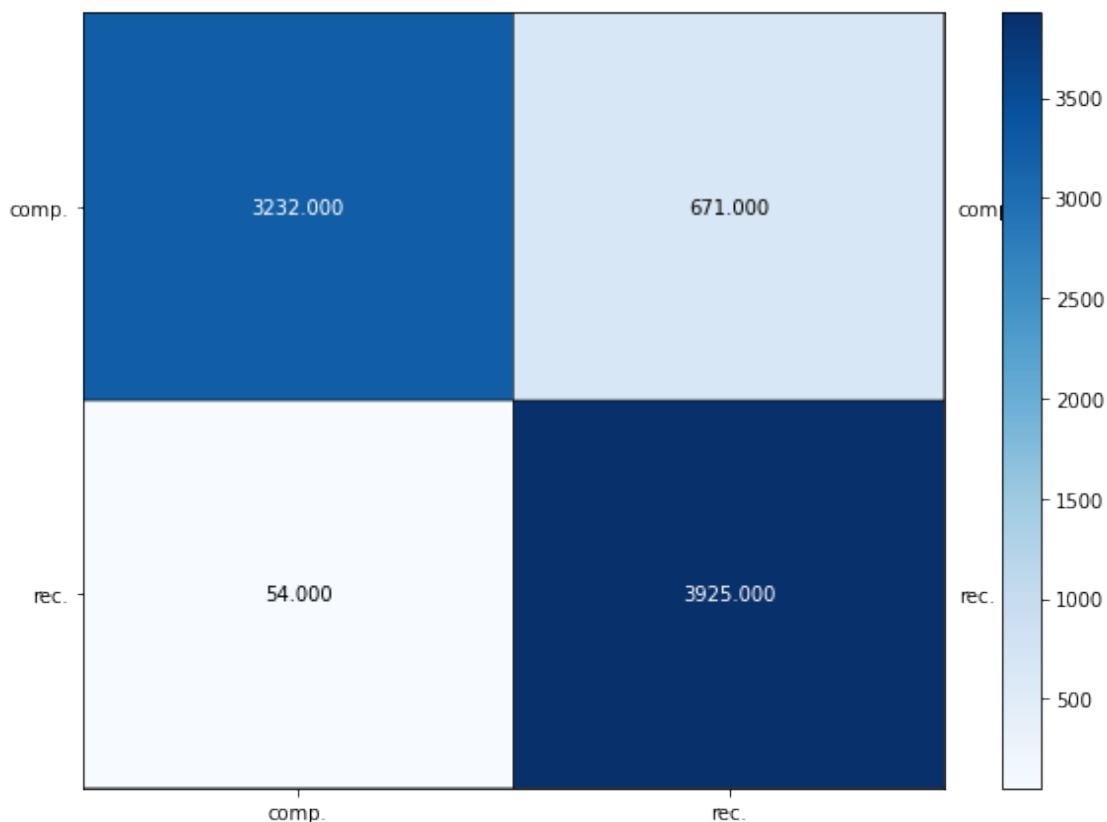
```
[ ]: # map_root = {"comp.graphics":0, "comp.os.ms-windows.misc":0, "comp.sys.ibm.pc.
↳ hardware":0, "comp.sys.mac.hardware":0,
```

```
#          "rec.autos":1, "rec.motorcycles":1, "rec.sport.baseball":1, "rec.
↳sport.hockey":1}
# y_data = dataset.target.map(map_root)
y_data = []
for i in dataset.target:
    if i < 4:
        y_data.append(0)
    else:
        y_data.append(1)
```

```
[ ]: km = KMeans(n_clusters=2, init='k-means++', max_iter=2000, n_init=50,↳
↳random_state=0)
km.fit(X_tfidf)
```

```
[ ]: KMeans(max_iter=2000, n_clusters=2, n_init=50, random_state=0)
```

```
[ ]: plot_mat(contingency_matrix(y_data, km.labels_), size=(8, 6),↳
↳xticklabels=['comp.', 'rec.'], yticklabels=['comp.', 'rec.'], pic_fname='Q2.
↳png')
```



0.4 Question 3

Report the 5 clustering measures explained in the introduction for Kmeans clustering.

Ans:

- Homogeneity: 0.5999
- Completeness: 0.6121
- V-measure: 0.6059
- Adjusted Rand-Index: 0.6659
- Adjusted Mutual Information Score: 0.6059

```
[ ]: def print_5_measure_scores(y_data, labels):  
    print("Homogeneity: %0.4f" % homogeneity_score(y_data, labels))  
    print("Completeness: %0.4f" % completeness_score(y_data, labels))  
    print("V-measure: %0.4f" % v_measure_score(y_data, labels))  
    print("Adjusted Rand-Index: %0.4f" % adjusted_rand_score(y_data, labels))  
    print("Adjusted Mutual Information Score: %0.4f"%  
    ↪adjusted_mutual_info_score(y_data, labels))  
print_5_measure_scores(y_data, km.labels_)
```


```
Homogeneity: 0.5999  
Completeness: 0.6121  
V-measure: 0.6059  
Adjusted Rand-Index: 0.6659  
Adjusted Mutual Information Score: 0.6059
```

0.5 Question 4

Report the plot of the percentage of variance that the top r principle components retain v.s. r, for r = 1 to 1000.

```
[ ]: svd = TruncatedSVD(n_components=1000, random_state=0)  
X_LSI = svd.fit_transform(X_tfidf)  
  
ratios = np.cumsum(svd.explained_variance_ratio_)  
fig = plt.figure()  
ax = fig.add_subplot(1, 1, 1)  
ax.plot(np.linspace(1, 1000, 1000), ratios, lw=2, linestyle='--')  
  
ax.set_ylabel('cumulative sum of explained_variance_ratio')  
ax.set_xlabel('principal components (r)')  
plt.title('the percentage of variance that the top r principle components_  
    ↪retain v.s. r')  
plt.show()
```

the percentage of variance that the top r principle components retain v.s. r



0.6 Question 5

Let r be the dimension that we want to reduce the data to (i.e. n components). Try $r = 1, 10, 20, 50, 100, 300$, and plot the 5 measure scores v.s. r for both SVD and NMF.

```
[ ]: rs5 = [1, 10, 20, 50, 100, 300]
# svd
svd_homo = []
svd_comp = []
svd_vmes = []
svd_ranI = []
svd_muts = []
# nmf
nmf_homo = []
nmf_comp = []
nmf_vmes = []
nmf_ranI = []
nmf_muts = []
kmm = KMeans(n_clusters=2, init='k-means++', max_iter=2000, n_init=50,
↳random_state=0)
for i in range(len(rs5)):
    print("Try r =", rs5[i], "...")
    # SVD
    svd_tmp = TruncatedSVD(n_components=rs5[i], random_state=0)
    X_svd = svd_tmp.fit_transform(X_tfidf)
    km_svd = kmm.fit(X_svd)
```



```

svd_homo.append(homogeneity_score(y_data, km_svd.labels_))
svd_comp.append(completeness_score(y_data, km_svd.labels_))
svd_vmes.append(v_measure_score(y_data, km_svd.labels_))
svd_ranI.append(adjusted_rand_score(y_data, km_svd.labels_))
svd_muts.append(adjusted_mutual_info_score(y_data, km_svd.labels_))
# NMF
nmf_tmp = NMF(n_components=rs5[i], init='random', random_state=0)
X_nmf = nmf_tmp.fit_transform(X_tfidf)
km_nmf = kmm.fit(X_nmf)
nmf_homo.append(homogeneity_score(y_data, km_nmf.labels_))
nmf_comp.append(completeness_score(y_data, km_nmf.labels_))
nmf_vmes.append(v_measure_score(y_data, km_nmf.labels_))
nmf_ranI.append(adjusted_rand_score(y_data, km_nmf.labels_))
nmf_muts.append(adjusted_mutual_info_score(y_data, km_nmf.labels_))
print('Done')

```

Try r = 1 ...

Try r = 10 ...

Try r = 20 ...

Try r = 50 ...

Try r = 100 ...

/usr/local/lib/python3.8/dist-packages/sklearn/decomposition/_nmf.py:1637:
ConvergenceWarning: Maximum number of iterations 200 reached. Increase it to
improve convergence.

```
warnings.warn(
```

Try r = 300 ...

Done

```

[ ]: def print_bar_scores_result(width, rs, svd_score, nmf_score, score_name, title):
    fig = plt.figure()
    ax = fig.add_subplot(1, 1, 1)
    # ax.plot(rs, svd_homo, color='r', label="SVD")
    # ax.plot(rs, nmf_homo, color='b', label="NMF")
    ax.bar(np.arange(len(rs)) - width/2, svd_score, width, label="SVD")
    ax.bar(np.arange(len(rs)) + width/2, nmf_score, width, label="NMF")
    ax.set_xticks(np.arange(len(rs)))
    ax.set_xticklabels(rs)

    ax.legend()
    ax.set_ylabel(score_name)
    ax.set_xlabel('r')
    plt.title(title)
    plt.show()

```

```


[ ]: print_bar_scores_result(0.3, rs5, svd_homo, nmf_homo, 'Homogeneity',
    ↪ 'Homogeneity Score for different r [SVD vs. NMF]')


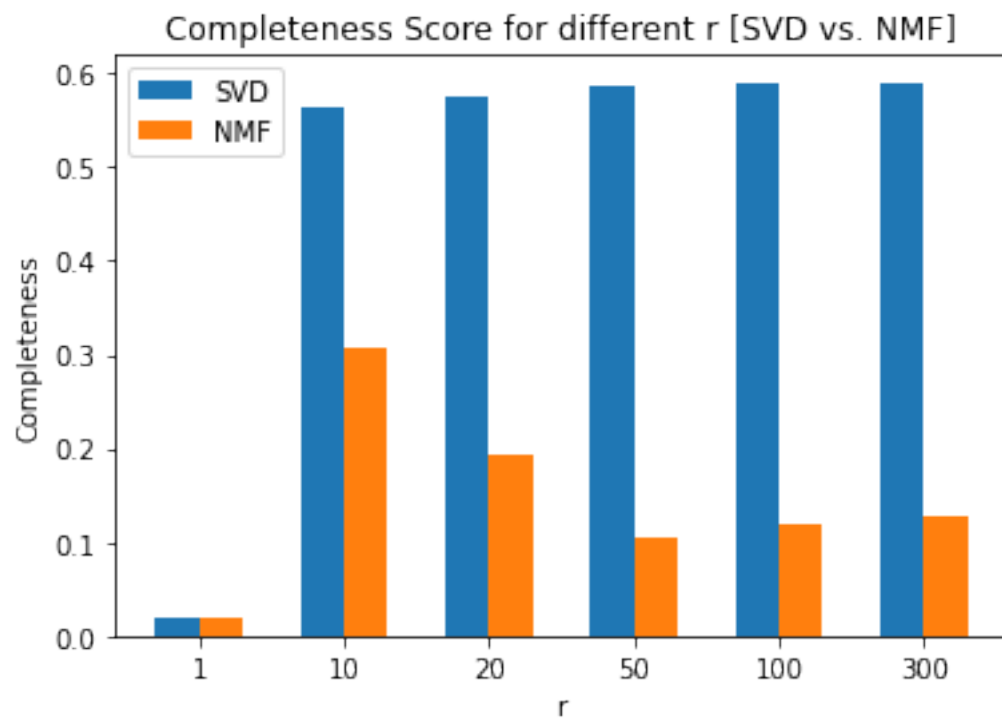
```



```

print_bar_scores_result(0.3, rs5, svd_comp, nmf_comp, 'Completeness',
    ↳ 'Completeness Score for different r [SVD vs. NMF]')
print_bar_scores_result(0.3, rs5, svd_vmes, nmf_vmes, 'V-measure', 'V-measure',
    ↳ 'Score for different r [SVD vs. NMF]')
print_bar_scores_result(0.3, rs5, svd_ranI, nmf_ranI, 'Adjusted Rand-Index',
    ↳ 'Adjusted Rand-Index for different r [SVD vs. NMF]')
print_bar_scores_result(0.3, rs5, svd_muts, nmf_muts, 'Adjusted Mutual',
    ↳ 'Information Score', 'Adjusted Mutual Information Score for different r [SVD',
    ↳ 'vs. NMF]')

```







Report a good choice of r for SVD and NMF respectively.

Ans:

From the graph above, we can find that a good choice of r for SVD could be $r = 50$, a good choice of r for NMF could be $r = 10$.

Because we can find that there are no significant increase for SVD after $r = 50$. And as r increases, we need to realize that the Euclidean distances in K-means we use will converge to a constant value, such that K-means doesn't perform well in high-dimensional prediction. Considering to this, choosing $r = 50$ is good enough. On the other hand, NMF performs obviously best when $r = 10$.

0.7 Question 6

How do you explain the non-monotonic behavior of the measures as r increases?

Ans:

As r increases, the performance of the measures doesn't increase correspondingly. It's because that as r increases, it indicates that the target matrix behind it will become complicated and high-dimensional. In which restrained the performance of K-means as we mentioned before. The Euclidean distances in K-means will converge to a constant value between the sample points. The clustering method can't find centroids for different sample points so that its performance becomes poor.

On the other hand, we can find that SVD maintain the similar score performance even after r keeps increasing, however, NMF doesn't perform well after r is larger than 10. It might because of the fact that SVD is a much more deterministic method than NMF. As we discussed in Project 1, SVD is a more insightful method and is able to interpret high-dimensional data rather than NMF did. Since NMF only uses the positive entries in the reduced matrix and makes assumption about the missing values. SVD doesn't assume anything about the value. Therefore, it's less restricted and performs better in high-dimensional data.

0.8 Question 7

Are these measures on average better than those computed in Question 3?

Ans:

No, even though we minus the outlier data ($r = 1$), the average performance of these measures is still worse than those computed in Question 3.

```
[ ]: print("Measure scores computed in Question 3:")
      print_5_measure_scores(y_data, km.labels_)

      print("\nAverage measure scores for SVD:")
      print("Homogeneity: %0.4f" % (sum(svd_homo) / len(svd_homo)))
      print("Completeness: %0.4f" % (sum(svd_comp) / len(svd_comp)))
      print("V-measure: %0.4f" % (sum(svd_vmes) / len(svd_vmes)))
      print("Adjusted Rand-Index: %0.4f" % (sum(svd_ranI) / len(svd_ranI)))
      print("Adjusted Mutual Information Score: %0.4f" % (sum(svd_muts) /
      ↪len(svd_muts)))
```

```

print("\nAverage measure scores for NMF:")
print("Homogeneity: %0.4f" % (sum(nmf_homo) / len(nmf_homo)))
print("Completeness: %0.4f" % (sum(nmf_comp) / len(nmf_comp)))
print("V-measure: %0.4f" % (sum(nmf_vmes) / len(nmf_vmes)))
print("Adjusted Rand-Index: %0.4f" % (sum(nmf_ranI) / len(nmf_ranI)))
print("Adjusted Mutual Information Score: %0.4f" % (sum(nmf_muts) /
    ↳len(nmf_muts)))

print("\nAverage measure scores for SVD (not including r=1):")
print("Homogeneity: %0.4f" % ((sum(svd_homo) - svd_homo[0]) / (len(svd_homo) -
    ↳1)))
print("Completeness: %0.4f" % ((sum(svd_comp) - svd_comp[0]) / (len(svd_comp) -
    ↳1)))
print("V-measure: %0.4f" % ((sum(svd_vmes) - svd_vmes[0]) / (len(svd_vmes) -
    ↳1)))
print("Adjusted Rand-Index: %0.4f" % ((sum(svd_ranI) - svd_ranI[0]) /
    ↳(len(svd_ranI) - 1)))
print("Adjusted Mutual Information Score: %0.4f" % ((sum(svd_muts) -
    ↳svd_ranI[0]) / (len(svd_muts) - 1)))

```

Measure scores computed in Question 3:

Homogeneity: 0.5999
 Completeness: 0.6121
 V-measure: 0.6059
 Adjusted Rand-Index: 0.6659
 Adjusted Mutual Information Score: 0.6059

Average measure scores for SVD:

Homogeneity: 0.4746
 Completeness: 0.4875
 V-measure: 0.4810
 Adjusted Rand-Index: 0.5254
 Adjusted Mutual Information Score: 0.4809

Average measure scores for NMF:

Homogeneity: 0.0528
 Completeness: 0.1453
 V-measure: 0.0680
 Adjusted Rand-Index: 0.0307
 Adjusted Mutual Information Score: 0.0679

Average measure scores for SVD (not including r=1):

Homogeneity: 0.5658
 Completeness: 0.5811
 V-measure: 0.5733
 Adjusted Rand-Index: 0.6253

Adjusted Mutual Information Score: 0.5719


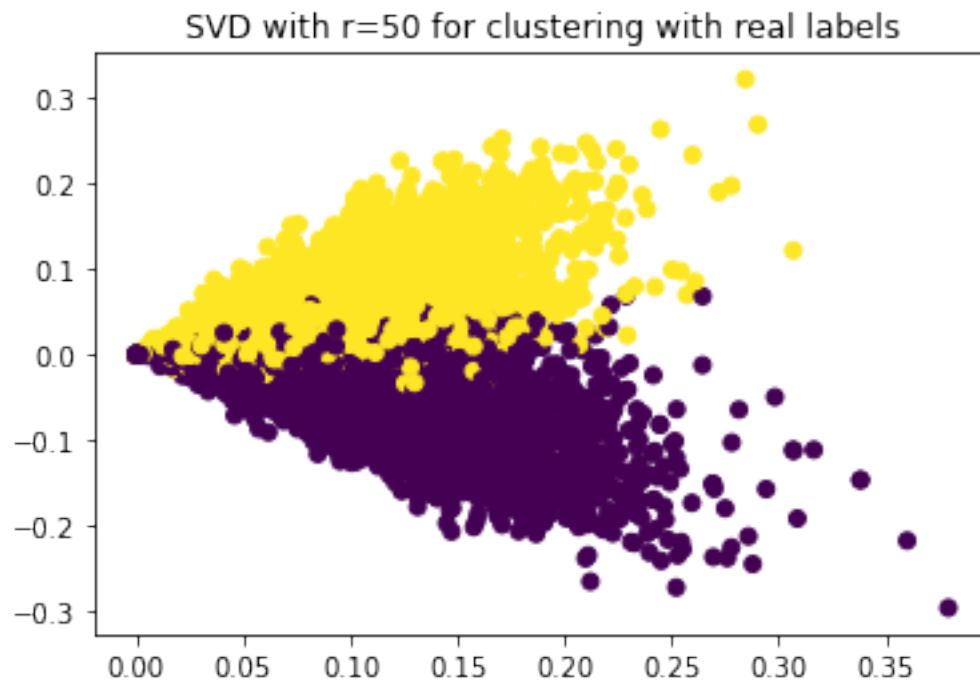
0.9 Question 8

Visualize the clustering results for:

- SVD with your optimal choice of r for K-Means clustering;
- NMF with your choice of r for K-Means clustering.

```
[ ]: # best parameter of SVD and NMF
best_r_SVD = 50
best_r_NMF = 10
kmm = KMeans(n_clusters=2, init='k-means++', max_iter=2000, n_init=50,
             random_state=0)
# best SVD
best_svd = TruncatedSVD(n_components=best_r_SVD, random_state=0)
best_X_svd = best_svd.fit_transform(X_tfidf)
y_pred_best_svd = kmm.fit_predict(best_X_svd)
# best NMF
best_nmf = NMF(n_components=best_r_NMF, init='random', random_state=0)
best_X_nmf = best_nmf.fit_transform(X_tfidf)
y_pred_best_nmf = kmm.fit_predict(best_X_nmf)
```


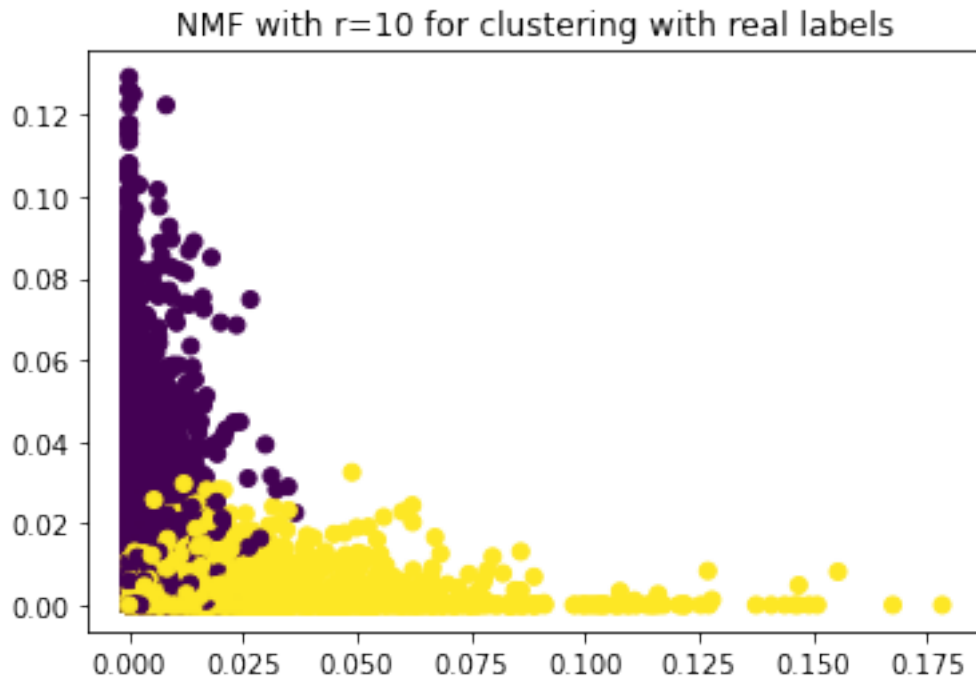
```
[ ]: plt.scatter(best_X_svd[:,0], best_X_svd[:,1], c=y_data)
plt.title("SVD with r=50 for clustering with real labels")
plt.show()
plt.scatter(best_X_svd[:,0], best_X_svd[:,1], c=y_pred_best_svd)
plt.title("SVD with r=50 for clustering with K-means predicted labels")
plt.show()
```



```
[ ]: plt.scatter(best_X_nmf[:,0], best_X_nmf[:,1], c=y_data)
plt.title("NMF with  $r=10$  for clustering with real labels")
```



```
plt.show()
plt.scatter(best_X_nmf[:,0], best_X_nmf[:,1], c=y_pred_best_nmf)
plt.title("NMF with r=10 for clustering with K-means predicted labels")
plt.show()
```



0.10 Question 9

What do you observe in the visualization? How are the data points of the two classes distributed? Is distribution of the data ideal for K-Means clustering?

Ans:

1. From the visualization, we can see that there are no spherical distributions on the plots for both SVD or NMF. there are many overlapping points among the two clusters which makes the boundary of two clusters hard to define.
2. The data points in the SVD are distributed more ideally than NMF did. Its shape also looks more like a sphere comparing to NMF one. On the other hand, we can analyze the performances from their homogeneity scores. SVD gets a higher score than NMF did but it's still not great enough.
3. In conclusion, the distribution of the data is not ideal for K-Means clustering. K-Means clustering assumes spherical shapes of clusters but we can not find a similar shape no matter in SVD or NMF here.

0.11 Question 10

Load documents with the same configuration as in Question 1, but for ALL 20 categories.

There is a mismatch between cluster labels and class labels. For example, the cluster #3 may correspond to the class #8. As a result, the high-value entries of the 20×20 contingency matrix can be scattered around, making it messy to inspect, even if the clustering result is not bad.

```
[ ]: dataset_all = fetch_20newsgroups(subset = 'all', shuffle = True, random_state = 0,
    ↪ remove=('headers', 'footers'))

vec = CountVectorizer(stop_words='english', min_df=3)
tfidf = TfidfTransformer()

X_a_vec = vec.fit_transform(dataset_all.data)
X_a_tfidf = tfidf.fit_transform(X_a_vec)
print("TFIDF dataset shape:", X_a_tfidf.shape)
```

TFIDF dataset shape: (18846, 45365)

```
[ ]: y_a_data = dataset_all.target

kma = KMeans(n_clusters=20, init='k-means++', max_iter=2000, n_init=50,
    ↪ random_state=0)
# kma.fit(X_a_tfidf)
```

Construct the TF-IDF matrix, reduce its dimensionality using BOTH NMF and SVD (specify settings you choose and why).

Ans:

We choose $r = 20$ for SVD, $r = 10$ for NMF respectively. we will explain the reason in the following cells.

```
[ ]: rs10 = [1, 5, 10, 20, 50, 100]
# svd
a_svd_homo = []
a_svd_comp = []
a_svd_vmes = []
a_svd_ranI = []
a_svd_muts = []
# nmf
a_nmf_homo = []
a_nmf_comp = []
a_nmf_vmes = []
a_nmf_ranI = []
a_nmf_muts = []
for i in range(len(rs10)):
    print("Try r =", rs10[i], "...")
    # SVD
    svd_tmp = TruncatedSVD(n_components=rs10[i], random_state=0)
    X_a_svd = svd_tmp.fit_transform(X_a_tfidf)
    kma_svd = kma.fit(X_a_svd)
    a_svd_homo.append(homogeneity_score(y_a_data, kma_svd.labels_))
    a_svd_comp.append(completeness_score(y_a_data, kma_svd.labels_))
    a_svd_vmes.append(v_measure_score(y_a_data, kma_svd.labels_))
    a_svd_ranI.append(adjusted_rand_score(y_a_data, kma_svd.labels_))
    a_svd_muts.append(adjusted_mutual_info_score(y_a_data, kma_svd.labels_))
    # NMF
    nmf_tmp = NMF(n_components=rs10[i], init='random', random_state=0)
    X_a_nmf = nmf_tmp.fit_transform(X_a_tfidf)
    kma_nmf = kma.fit(X_a_nmf)
    a_nmf_homo.append(homogeneity_score(y_a_data, kma_nmf.labels_))
    a_nmf_comp.append(completeness_score(y_a_data, kma_nmf.labels_))
    a_nmf_vmes.append(v_measure_score(y_a_data, kma_nmf.labels_))
    a_nmf_ranI.append(adjusted_rand_score(y_a_data, kma_nmf.labels_))
    a_nmf_muts.append(adjusted_mutual_info_score(y_a_data, kma_nmf.labels_))
print('Done')
```

Try r = 1 ...

Try r = 5 ...

Try r = 10 ...

/usr/local/lib/python3.8/dist-packages/sklearn/decomposition/_nmf.py:1637:
ConvergenceWarning: Maximum number of iterations 200 reached. Increase it to
improve convergence.


warnings.warn(

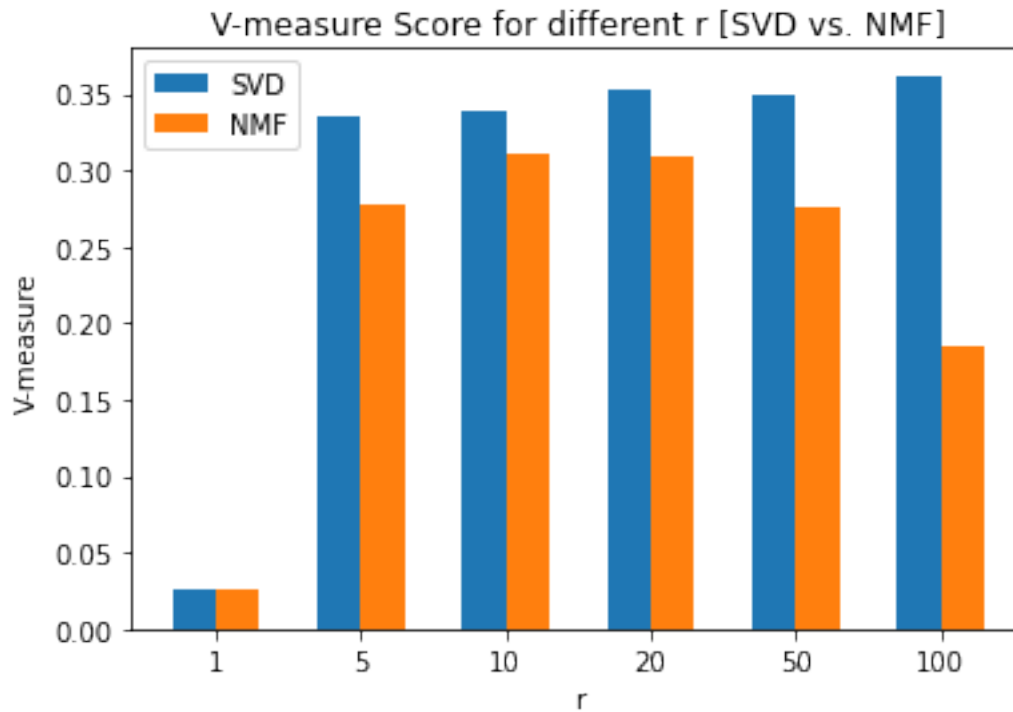
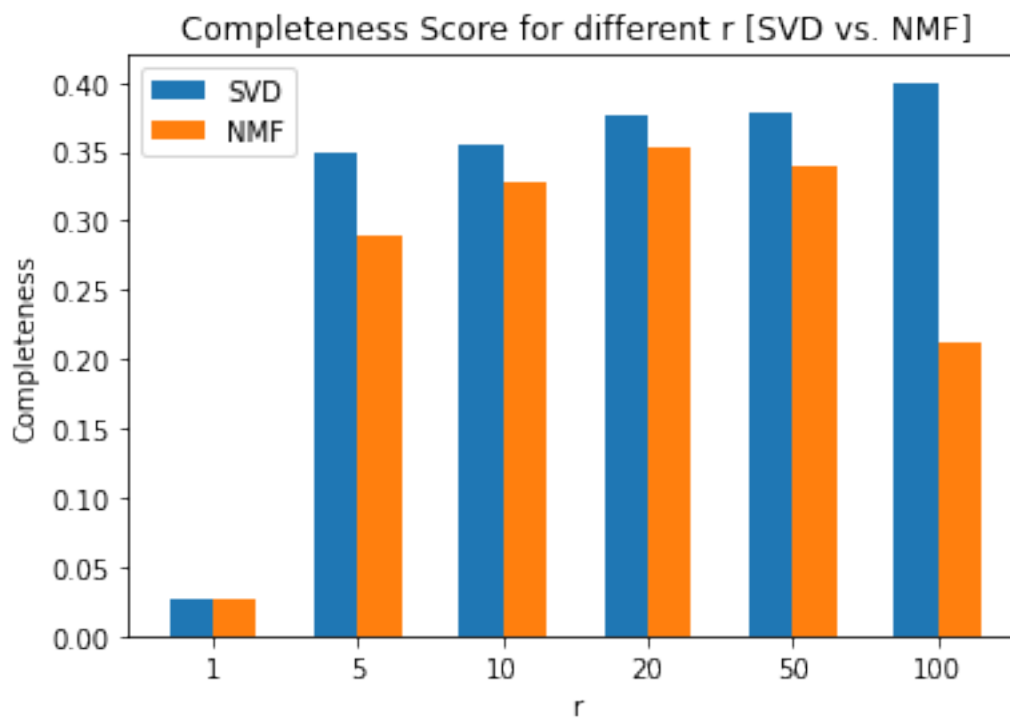
Try r = 20 ...


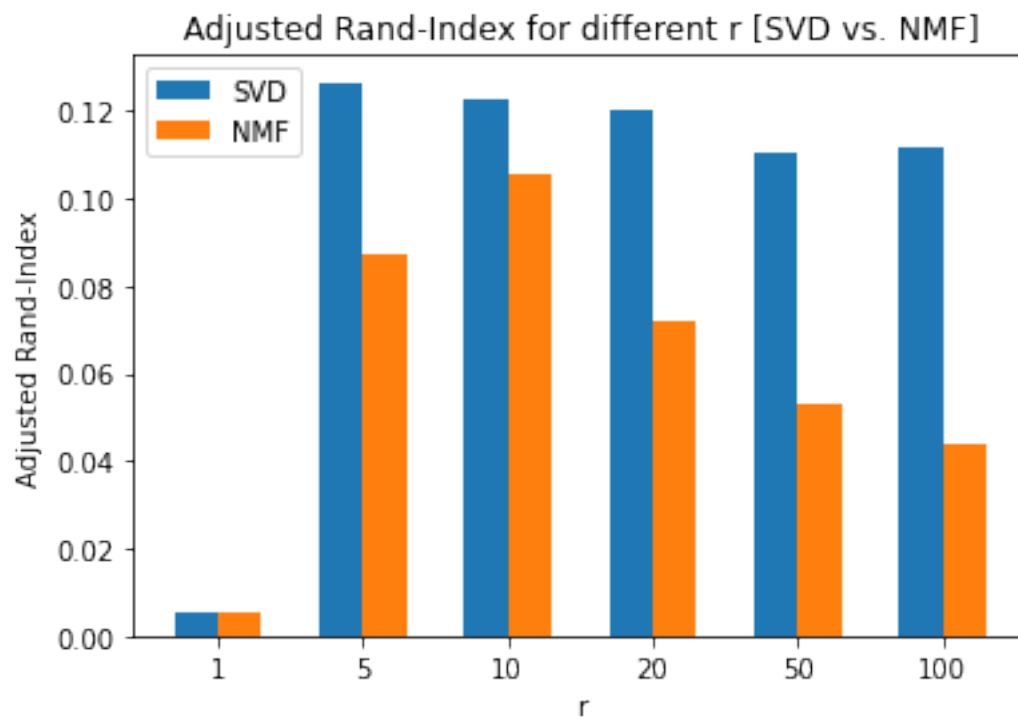
Try r = 50 ...

Try r = 100 ...
Done

```
[ ]: print_bar_scores_result(0.3, rs10, a_svd_homo, a_nmf_homo, 'Homogeneity',  
    ↳ 'Homogeneity Score for different r [SVD vs. NMF]')  
print_bar_scores_result(0.3, rs10, a_svd_comp, a_nmf_comp, 'Completeness',  
    ↳ 'Completeness Score for different r [SVD vs. NMF]')  
print_bar_scores_result(0.3, rs10, a_svd_vmes, a_nmf_vmes, 'V-measure',  
    ↳ 'V-measure Score for different r [SVD vs. NMF]')  
print_bar_scores_result(0.3, rs10, a_svd_ranI, a_nmf_ranI, 'Adjusted_  
    ↳ Rand-Index', 'Adjusted Rand-Index for different r [SVD vs. NMF]')  
print_bar_scores_result(0.3, rs10, a_svd_muts, a_nmf_muts, 'Adjusted Mutual_  
    ↳ Information Score', 'Adjusted Mutual Information Score for different r [SVD_  
    ↳ vs. NMF]')
```







```

[ ]: avg_svd_homo = sum(a_svd_homo) / len(a_svd_homo)
avg_svd_comp = sum(a_svd_comp) / len(a_svd_comp)
avg_svd_vmes = sum(a_svd_vmes) / len(a_svd_vmes)
avg_svd_ranI = sum(a_svd_ranI) / len(a_svd_ranI)
avg_svd_muts = sum(a_svd_muts) / len(a_svd_muts)

avg_nmf_homo = sum(a_nmf_homo) / len(a_nmf_homo)
avg_nmf_comp = sum(a_nmf_comp) / len(a_nmf_comp)
avg_nmf_vmes = sum(a_nmf_vmes) / len(a_nmf_vmes)
avg_nmf_ranI = sum(a_nmf_ranI) / len(a_nmf_ranI)
avg_nmf_muts = sum(a_nmf_muts) / len(a_nmf_muts)

avg_svd = []
avg_nmf = []
for i in range(len(rs10)):
    tmp_svd = 0
    tmp_svd += a_svd_homo[i] / avg_svd_homo
    tmp_svd += a_svd_comp[i] / avg_svd_comp
    tmp_svd += a_svd_vmes[i] / avg_svd_vmes
    tmp_svd += a_svd_ranI[i] / avg_svd_ranI
    tmp_svd += a_svd_muts[i] / avg_svd_muts
    avg_svd.append(tmp_svd / 5)
    tmp_nmf = 0
    tmp_nmf += a_nmf_homo[i] / avg_nmf_homo
    tmp_nmf += a_nmf_comp[i] / avg_nmf_comp
    tmp_nmf += a_nmf_vmes[i] / avg_nmf_vmes
    tmp_nmf += a_nmf_ranI[i] / avg_nmf_ranI
    tmp_nmf += a_nmf_muts[i] / avg_nmf_muts
    avg_nmf.append(tmp_nmf / 5)
for i in range(len(avg_svd)):
    print("Average normalized scores for SVD when r =", rs10[i], ":", avg_svd[i])
print("")
for i in range(len(avg_nmf)):
    print("Average normalized scores for NMF when r =", rs10[i], ":", avg_nmf[i])

sorted_avg_svd = sorted(avg_svd, reverse=True)
sorted_avg_nmf = sorted(avg_nmf, reverse=True)

top3_svd = [avg_svd.index(v) for v in sorted_avg_svd[:3]]
top3_nmf = [avg_nmf.index(v) for v in sorted_avg_nmf[:3]]

print("\nTop 3 best values of r for SVD:")
for i in top3_svd:
    print("r =", rs10[i], ":", avg_svd[i])

print("\nTop 3 best values of r for NMF:")
for i in top3_nmf:

```

```
print("r =", rs10[i], ":", avg_nmf[i])
```

```
Average normalized scores for SVD when r = 1 : 0.07766046505240246
Average normalized scores for SVD when r = 5 : 1.164613083355113
Average normalized scores for SVD when r = 10 : 1.169345407002839
Average normalized scores for SVD when r = 20 : 1.2027367831942999
Average normalized scores for SVD when r = 50 : 1.174476757975515
Average normalized scores for SVD when r = 100 : 1.2111675034198306
```

```
Average normalized scores for NMF when r = 1 : 0.10256664314454249
Average normalized scores for NMF when r = 5 : 1.2478007770753372
Average normalized scores for NMF when r = 10 : 1.4200459030547532
Average normalized scores for NMF when r = 20 : 1.3052363009210974
Average normalized scores for NMF when r = 50 : 1.1389088770343618
Average normalized scores for NMF when r = 100 : 0.7854414987699082
```

Top 3 best values of r for SVD:

```
r = 100 : 1.2111675034198306
r = 20 : 1.2027367831942999
r = 50 : 1.174476757975515
```

Top 3 best values of r for NMF:

```
r = 10 : 1.4200459030547532
r = 20 : 1.3052363009210974
r = 5 : 1.2478007770753372
```

We computed the 5 measure scores, normalized it to sum up and get average. As you can see, we obtained the top 3 best values of r for SVD and NMF respectively.


Based on above result, we choose $r = 20$ for SVD, $r = 10$ for NMF as setting.

As r increases to 100 for SVD, there is no corresponding significant increasing in the normalized scores for it. Therefore, choosing $r = 20$ is good enough.

Choosing $r = 10$ for NMF is easy here, since it's lower than 20 but get a higher normalized average score.

Visualize the contingency matrix and report the five clustering metrics (DO BOTH NMF AND SVD).


```
[ ]: best_a_svd = 20
# SVD
best_svd = TruncatedSVD(n_components=best_a_svd, random_state=0)
best_X_a_svd = best_svd.fit_transform(X_a_tfidf)
best_kma_svd = kma.fit(best_X_a_svd)
cm = confusion_matrix(dataset_all.target, kma.labels_)
rows, cols = linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
↪title='SVD with best r (r = 20)', size=(15,15))
```

```
[ ]: best_a_nmf = 10
# NMF
best_nmf = NMF(n_components=best_a_nmf, random_state=0)
best_X_a_nmf = best_nmf.fit_transform(X_a_tfidf)
best_kma_nmf = kma.fit(best_X_a_nmf)
cm = confusion_matrix(dataset_all.target, kma.labels_)
rows, cols = linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], cols, xticklabels=cols, yticklabels=rows,
↪title='NMF with best r (r = 10)', size=(15,15))
```

/usr/local/lib/python3.8/dist-packages/sklearn/decomposition/_nmf.py:289:
FutureWarning: The 'init' value, when 'init=None' and n_components is less than
n_samples and n_features, will be changed from 'nndsvd' to 'nndsvda' in 1.1

```
(renaming of 0.26).
warnings.warn(
```



0.12 Question 11

Reduce the dimension of your dataset with UMAP. Consider the following settings: n components = [5, 20, 200], metric = "cosine" vs. "euclidean". If "cosine" metric fails, please look at the FAQ at the end of this spec.

Report the permuted contingency matrix and the five clustering evaluation metrics for the different combinations (6 combinations).

```
[ ]: rs11 = [5, 20, 200]
      # metrics = ['cosine', 'euclidean']
```

```

# UMAP cosine
umap_cos_homo = []
umap_cos_comp = []
umap_cos_vmes = []
umap_cos_ranI = []
umap_cos_muts = []
# UMAP euclidean
umap_euc_homo = []
umap_euc_comp = []
umap_euc_vmes = []
umap_euc_ranI = []
umap_euc_muts = []

kma = KMeans(n_clusters=20, init='k-means++', max_iter=2000, n_init=50,
↳random_state=0)

for i in range(len(rs11)):
    print("\nTry r =", rs11[i], "...")
    # cosine
    print("Cosine UMAP when r =", rs11[i], ":")
    tmp_umap_cos = umap.UMAP(n_components=rs11[i], metric='cosine')
    X_umap_cos = tmp_umap_cos.fit_transform(X_a_tfidf)
    kma_cos = kma.fit(X_umap_cos)
    umap_cos_homo.append(homogeneity_score(y_a_data, kma_cos.labels_))
    umap_cos_comp.append(completeness_score(y_a_data, kma_cos.labels_))
    umap_cos_vmes.append(v_measure_score(y_a_data, kma_cos.labels_))
    umap_cos_ranI.append(adjusted_rand_score(y_a_data, kma_cos.labels_))
    umap_cos_muts.append(adjusted_mutual_info_score(y_a_data, kma_cos.labels_))
    # print permuted contingency matrix
    cm = confusion_matrix(y_a_data, kma.labels_)
    rows, cols = linear_sum_assignment(cm, maximize=True)
    plot_mat(cm[rows[:, np.newaxis], cols], cols, xticklabels=cols, yticklabels=rows,
↳title='cosine UMAP with r=' + str(rs11[i]), size=(15,15))
    # print the five clustering evaluation metrics
    print_5_measure_scores(y_a_data, kma_cos.labels_)
    # euclidean
    print("\nEuclidean UMAP when r =", rs11[i], ":")
    tmp_umap_euc = umap.UMAP(n_components=rs11[i], metric='euclidean')
    X_umap_euc = tmp_umap_cos.fit_transform(X_a_tfidf)
    kma_euc = kma.fit(X_umap_euc)
    umap_euc_homo.append(homogeneity_score(y_a_data, kma_euc.labels_))
    umap_euc_comp.append(completeness_score(y_a_data, kma_euc.labels_))
    umap_euc_vmes.append(v_measure_score(y_a_data, kma_euc.labels_))
    umap_euc_ranI.append(adjusted_rand_score(y_a_data, kma_euc.labels_))
    umap_euc_muts.append(adjusted_mutual_info_score(y_a_data, kma_euc.labels_))
    # print permuted contingency matrix
    cm = confusion_matrix(y_a_data, kma.labels_)

```


```

rows, cols = linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
→title='euclidean UMAP with r=' + str(rs11[i]), size=(15,15))
# print the five clustering evaluation metrics
print_5_measure_scores(y_a_data, kma_euc.labels_)
print('Done')

```

Try r = 5 ...

Cosine UMAP when r = 5 :




Homogeneity: 0.5566

Completeness: 0.5911

V-measure: 0.5733
Adjusted Rand-Index: 0.4403
Adjusted Mutual Information Score: 0.5719


Euclidean UMAP when $r = 5$:



Homogeneity: 0.5651
Completeness: 0.5910
V-measure: 0.5778
Adjusted Rand-Index: 0.4451
Adjusted Mutual Information Score: 0.5764


Try $r = 20$...

Cosine UMAP when $r = 20$:




Homogeneity: 0.5709
 Completeness: 0.5962
 V-measure: 0.5833
 Adjusted Rand-Index: 0.4623
 Adjusted Mutual Information Score: 0.5819

Euclidean UMAP when r = 20 :




Homogeneity: 0.5691
 Completeness: 0.5957
 V-measure: 0.5821
 Adjusted Rand-Index: 0.4443
 Adjusted Mutual Information Score: 0.5807

Try r = 200 ...
 Cosine UMAP when r = 200 :



Homogeneity: 0.5623
 Completeness: 0.5895
 V-measure: 0.5756
 Adjusted Rand-Index: 0.4471
 Adjusted Mutual Information Score: 0.5742

Euclidean UMAP when r = 200 :



Homogeneity: 0.5852
Completeness: 0.6009
V-measure: 0.5930
Adjusted Rand-Index: 0.4623
Adjusted Mutual Information Score: 0.5916
Done

0.13 Question 12

Analyze the contingency matrices. Which setting works best and why?

Ans:

From the contingency metrices we found above, we can see that there are no many differences between the euclidean UMAP with different r . Therefore, setting $r = 20$ is good enough for

euclidean UMAP.

As of cosine UMAP, setting $r = 20$ gives us a better result. Setting r equals to other numbers all gives more outliers or categories being distributed to other categories.

```
[ ]: avg_cos = []
avg_euc = []
for i in range(len(rs11)):
    tmp_cos = 0
    tmp_cos += umap_cos_homo[i] / (sum(umap_cos_homo) / len(umap_cos_homo))
    tmp_cos += umap_cos_comp[i] / (sum(umap_cos_comp) / len(umap_cos_comp))
    tmp_cos += umap_cos_vmes[i] / (sum(umap_cos_vmes) / len(umap_cos_vmes))
    tmp_cos += umap_cos_ranI[i] / (sum(umap_cos_ranI) / len(umap_cos_ranI))
    tmp_cos += umap_cos_muts[i] / (sum(umap_cos_muts) / len(umap_cos_muts))
    avg_cos.append(tmp_cos / 5)
    tmp_euc = 0
    tmp_euc += umap_euc_homo[i] / (sum(umap_euc_homo) / len(umap_euc_homo))
    tmp_euc += umap_euc_comp[i] / (sum(umap_euc_comp) / len(umap_euc_comp))
    tmp_euc += umap_euc_vmes[i] / (sum(umap_euc_vmes) / len(umap_euc_vmes))
    tmp_euc += umap_euc_ranI[i] / (sum(umap_euc_ranI) / len(umap_euc_ranI))
    tmp_euc += umap_euc_muts[i] / (sum(umap_euc_muts) / len(umap_euc_muts))
    avg_euc.append(tmp_euc / 5)
for i in range(len(avg_cos)):
    print("Average normalized scores for cosine UMAP when r =", rs11[i], ":", "\n"
    → avg_cos[i])
print("")
for i in range(len(avg_euc)):
    print("Average normalized scores for euclidean UMAP when r =", rs11[i], ":", "\n"
    → avg_euc[i])

best_r_cos = rs11[avg_cos.index(max(avg_cos))]
# best_r_cos = 5
best_r_euc = rs11[avg_euc.index(max(avg_euc))]
print("\nBest value of r for cosine UMAP:", best_r_cos, ", its average \n"
    → normalized scores:", avg_cos[rs11.index(best_r_cos)])
print("Best value of r for euclidean UMAP:", best_r_euc, ", its average \n"
    → normalized scores:", avg_euc[rs11.index(best_r_euc)])
```

Average normalized scores for cosine UMAP when $r = 5$: 0.9901320195164086
Average normalized scores for cosine UMAP when $r = 20$: 1.0136027226843964
Average normalized scores for cosine UMAP when $r = 200$: 0.9962652577991953

Average normalized scores for euclidean UMAP when $r = 5$: 0.9886982164174374
Average normalized scores for euclidean UMAP when $r = 20$: 0.9942207323171074
Average normalized scores for euclidean UMAP when $r = 200$: 1.0170810512654551

Best value of r for cosine UMAP: 20 , its average normalized scores:
1.0136027226843964

Best value of r for euclidean UMAP: 200 , its average normalized scores:
1.0170810512654551

What about for each metric choice?

Ans:

From the result above, we calculated the average normalized scores for each combination.

It also indicates that there are no many differences between the euclidean UMAP with different r . Although the scores when $r = 200$ is a little bit better, choosing $r = 20$ is good enough considering to the calculation time.

On the other hand, a good choice of r for cosine UMAP is 20.

Comparing the scores of these two UMAP, cosine UMAP with $r = 20$ works best and we choose it for the following comparison.

0.14 Question 13

So far, we have attempted K-Means clustering with 4 different representation learning techniques (sparse TF-IDF representation, PCA-reduced, NMF-reduced, UMAP-reduced).

Compare and contrast the clustering results across the 4 choices, and suggest an approach that is best for the K-Means clustering task on the 20-class text data. Choose any choice of clustering metrics for your comparison.

Ans:

Comparin the clustering results across the 4 choices, we can see that the performance of UMAP-reduced is better than other 3 results. Therefore, we suggest **UMAP-reduced (cosine UMAP with $r = 20$)** is best for the K-Means clustering task on the 20-class text data.

```
[ ]: def print_scores(title_name, score_homo, score_comp, score_vmes, score_ranI, \
    ↪score_muts):
    print(title_name)
    print("Homogeneity: %0.4f" % score_homo)
    print("Completeness: %0.4f" % score_comp)
    print("V-measure: %0.4f" % score_vmes)
    print("Adjusted Rand-Index: %0.4f" % score_ranI)
    print("Adjusted Mutual Information Score: %0.4f \n" % score_muts)
```

```
[ ]: kmm_o = KMeans(n_clusters=20, init='k-means++', max_iter=2000, n_init=50, \
    ↪random_state=0)
kmm_o.fit(X_a_tfidf)
```

```
[ ]: print("scores for sparse TF-IDF representation:")
print_5_measure_scores(y_a_data, kmm_o.labels_)
print("")
br_svd_i = rs10.index(best_a_svd)
br_nmf_i = rs10.index(best_a_nmf)
br_cos_i = avg_cos.index(max(avg_cos))
```

```

print_scores("scores for SVD when r = " + str(best_a_svd) + ":",
    ↪a_svd_homo[br_svd_i], a_svd_comp[br_svd_i], a_svd_vmes[br_svd_i],
    ↪a_svd_ranI[br_svd_i], a_svd_muts[br_svd_i])
print_scores("scores for NMF when r = " + str(best_a_nmf) + ":",
    ↪a_nmf_homo[br_nmf_i], a_nmf_comp[br_nmf_i], a_nmf_vmes[br_nmf_i],
    ↪a_nmf_ranI[br_nmf_i], a_nmf_muts[br_nmf_i])
print_scores("scores for cosine UMAP when r = " + str(rs11[br_cos_i]) + ":",
    ↪umap_cos_homo[br_cos_i], umap_cos_comp[br_cos_i], umap_cos_vmes[br_cos_i],
    ↪umap_cos_ranI[br_cos_i], umap_euc_muts[br_cos_i])

```

scores for sparse TF-IDF representation:

Homogeneity: 0.3470
 Completeness: 0.3928
 V-measure: 0.3684
 Adjusted Rand-Index: 0.1265
 Adjusted Mutual Information Score: 0.3663

scores for SVD when r = 20:

Homogeneity: 0.3330
 Completeness: 0.3756
 V-measure: 0.3530
 Adjusted Rand-Index: 0.1202
 Adjusted Mutual Information Score: 0.3508

scores for NMF when r = 10:

Homogeneity: 0.2953
 Completeness: 0.3275
 V-measure: 0.3106
 Adjusted Rand-Index: 0.1056
 Adjusted Mutual Information Score: 0.3082

scores for cosine UMAP when r = 20:

Homogeneity: 0.5709
 Completeness: 0.5962
 V-measure: 0.5833
 Adjusted Rand-Index: 0.4623
 Adjusted Mutual Information Score: 0.5807

0.15 Question 14

Use UMAP to reduce the dimensionality properly, and perform Agglomerative clustering with `n_clusters=20`. Compare the performance of “ward” and “single” linkage criteria.

Report the five clustering evaluation metrics for each case.

Ans:

From the five clustering evaluation metrics indicate below, we can find that the performance of

Agglomerative Clustering with single linkage criterion is much worse than ward linkage criterion one did.

It might be because of the using method behind these two criterions. The ward linkage criterion minimizes the variance of the clusters being merged. The single linkage criterion minimizes the distance between all observations of the two sets. Therefore, single linkage criterion is not robust enough to handle this high dimensional problem. It fails to deal with the noise or outlier data.

```
[ ]: best_umap = umap.UMAP(n_components=rs11[br_cos_i], metric='cosine')
X_umap = best_umap.fit_transform(X_a_tfidf)
agg_w = AgglomerativeClustering(n_clusters=20, linkage='ward').fit(X_umap)
agg_s = AgglomerativeClustering(n_clusters=20, linkage='single').fit(X_umap)

print("Agglomerative Clustering with ward linkage criterion:")
print_5_measure_scores(y_a_data, agg_w.labels_)
print("\nAgglomerative Clustering with single linkage criterion:")
print_5_measure_scores(y_a_data, agg_s.labels_)
```

Agglomerative Clustering with ward linkage criterion:

Homogeneity: 0.5594

Completeness: 0.5925

V-measure: 0.5755

Adjusted Rand-Index: 0.4295

Adjusted Mutual Information Score: 0.5741

Agglomerative Clustering with single linkage criterion:

Homogeneity: 0.0192

Completeness: 0.3767

V-measure: 0.0365

Adjusted Rand-Index: 0.0006

Adjusted Mutual Information Score: 0.0316

0.16 Question 15

Apply HDBSCAN on UMAP-transformed 20-category data.

Use min_cluster_size=100 .

Vary the min_cluster_size among 20, 100, 200 and report your findings in terms of the five clustering evaluation metrics - you will plot the best contingency matrix in the next question. Feel free to try modifying other parameters in HDBSCAN to get better performance.

Ans:

From the five clustering evaluation metrics we found, the best min_cluster_size to use here is 100.

```
[ ]: cluster_sizes = [20, 100, 200]
hdb_homo = []
hdb_comp = []
hdb_vmes = []
hdb_ranI = []
```

```

hdbb_muts = []

for min_size in cluster_sizes:
    print("\nmin_cluster_size =", min_size, ":")
    y_pred_hdbb = hdbbcan.HDBSCAN(min_cluster_size=min_size).fit_predict(X_umap)
    hdbb_homo.append(homogeneity_score(y_a_data, y_pred_hdbb))
    hdbb_comp.append(completeness_score(y_a_data, y_pred_hdbb))
    hdbb_vmes.append(v_measure_score(y_a_data, y_pred_hdbb))
    hdbb_ranI.append(adjusted_rand_score(y_a_data, y_pred_hdbb))
    hdbb_muts.append(adjusted_mutual_info_score(y_a_data, y_pred_hdbb))
    print_5_measure_scores(y_a_data, y_pred_hdbb)

```

```

min_cluster_size = 20 :
Homogeneity: 0.4253
Completeness: 0.4479
V-measure: 0.4363
Adjusted Rand-Index: 0.0777
Adjusted Mutual Information Score: 0.4236

```

```

min_cluster_size = 100 :
Homogeneity: 0.4157
Completeness: 0.6266
V-measure: 0.4998
Adjusted Rand-Index: 0.2263
Adjusted Mutual Information Score: 0.4989

```

```

min_cluster_size = 200 :
Homogeneity: 0.4200
Completeness: 0.6166
V-measure: 0.4997
Adjusted Rand-Index: 0.2191
Adjusted Mutual Information Score: 0.4987

```

```

[ ]: avg_hdbb = []
for i in range(len(cluster_sizes)):
    tmp_hdbb = 0
    tmp_hdbb += hdbb_homo[i] / (sum(hdbb_homo) / len(hdbb_homo))
    tmp_hdbb += hdbb_comp[i] / (sum(hdbb_comp) / len(hdbb_comp))
    tmp_hdbb += hdbb_vmes[i] / (sum(hdbb_vmes) / len(hdbb_vmes))
    tmp_hdbb += hdbb_ranI[i] / (sum(hdbb_ranI) / len(hdbb_ranI))
    tmp_hdbb += hdbb_muts[i] / (sum(hdbb_muts) / len(hdbb_muts))
    avg_hdbb.append(tmp_hdbb / 5)
for i in range(len(avg_hdbb)):
    print("Average normalized scores for HDBSCAN when min_cluster_size =",
    ↪cluster_sizes[i], ":", avg_hdbb[i])

```

```
best_mcs = cluster_sizes[avg_hdbs.index(max(avg_hdbs))]
print("\nBest value of min_cluster_size for HDBSCAN:", best_mcs, ", its average_
↳normalized scores:", max(avg_hdbs))
```

Average normalized scores for HDBSCAN when min_cluster_size = 20 :

0.8115684759351203

Average normalized scores for HDBSCAN when min_cluster_size = 100 :

1.0991266633883054

Average normalized scores for HDBSCAN when min_cluster_size = 200 :

1.0893048606765743

Best value of min_cluster_size for HDBSCAN: 100 , its average normalized scores:
1.0991266633883054

0.17 Question 16

Contingency matrix

Plot the contingency matrix for the best clustering model from Question 15. How many clusters are given by the model?

Interpret the contingency matrix considering the answer to these questions.

Ans:

According to our finding, there are total 10 major clusters given by the model. Furthermore, you can find there are many categories being distributed to other categories as a cluster.


It might be because of the reason that density based clustering relies on having enough data to separate dense areas. In higher dimensional spaces this becomes more difficult, and hence requires more data. Which makes HDBSCAN hard to perform well in this high dimensional data.

What does “-1” mean for the clustering labels?

Ans:

The “-1” in the clustering means noise or outlier samples that can not be classified into a cluster by the algorithms.

```
[ ]: # print contingency matrix
by_pred_hdbs = hdbscan.HDBSCAN(min_cluster_size=best_mcs).fit_predict(X_umap)
cm16 = confusion_matrix(y_a_data, by_pred_hdbs)
rows, cols = linear_sum_assignment(cm16, maximize=True)
plot_mat(cm16[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
↳title='HDBSCAN with min_cluster_size=' + str(best_mcs), size=(15,15))
```



0.18 Question 17

Based on your experiments, which dimensionality reduction technique and clustering methods worked best together for 20-class text data and why? Follow the table below. If UMAP takes too long to converge, consider running it once and saving the intermediate results in a pickle file.

| Module | Alternatives | Hyperparameters |
|--------------------------|--------------------------|--------------------------------|
| Dimensionality Reduction | None | N/A |
| Dimensionality Reduction | SVD | $r = [5, 20, 200]$ |
| Dimensionality Reduction | NMF | $r = [5, 20, 200]$ |
| Dimensionality Reduction | UMAP | $n_components = [5, 20, 200]$ |
| Clustering | K-Means | $k = [10, 20, 50]$ |
| Clustering | Agglomerative Clustering | $n_clusters = [20]$ |

| Module | Alternatives | Hyperparameters |
|------------|--------------|-------------------------------|
| Clustering | HDBSCAN | min_cluster_size = [100, 200] |

Ans:

The combination of **UMAP** using 'cosine' matrix and **n_components** setting to 5 plus **K-Means** using **k = 20** works best together for 20-class text data. we will analysis the reason of it in the following cells.

```
[68]: def pkl_save(path, data, filename):
        with open(path + filename, 'wb') as f:
            # compressed_file = bz2.BZ2File(f, 'w')
            pickle.dump(data, f)
```

```
[ ]: dataset_all = fetch_20newsgroups(subset = 'all', shuffle = True, random_state = 0,
    ↪ remove=('headers', 'footers'))
vec = CountVectorizer(stop_words='english', min_df=3)
tfidf = TfidfTransformer()
X_vec = vec.fit_transform(dataset_all.data)
X_tfidf = tfidf.fit_transform(X_vec)
y_data = dataset_all.target
```

```
[76]: # modify this line to your local path (you don't need this cell if you already
    ↪ specify your path in the first cell)
# path = '/content/drive/MyDrive/Colab Notebooks/ECE219/Project2/'
# path = '/Users/behind/Desktop/UCLA/Winter 2023/EC ENGR 219/Project2/pickle
    ↪ file/'
```

```
[ ]: # None
pkl_save(path, X_tfidf, 'none.pkl')

rs17 = [5, 20, 200]
for i in range(len(rs17)):
    print("\ntry r =", rs17[i], "...")
    # SVD
    print("using SVD ...")
    svd_tmp = TruncatedSVD(n_components=rs17[i], random_state=0)
    X_svd = svd_tmp.fit_transform(X_tfidf)
    pkl_save(path, X_svd, 'svd_' + str(rs17[i]) + '.pkl')

    # NMF
    print("using NMF ...")
    nmf_tmp = NMF(n_components=rs17[i], init='random', random_state=0)
    X_nmf = nmf_tmp.fit_transform(X_tfidf)
    pkl_save(path, X_nmf, 'nmf_' + str(rs17[i]) + '.pkl')

    # UMAP
```

```

print("using UMAP ...")
umap_tmp = umap.UMAP(n_components=rs17[i], metric='cosine')
X_umap = umap_tmp.fit_transform(X_tfidf)
pkl_save(path, X_umap, 'umap_' + str(rs17[i]) + '.pkl')

```

```

try r = 5 ...
using SVD ...
using NMF ...
using UMAP ...

```

```

try r = 20 ...
using SVD ...
using NMF ...
using UMAP ...

```

```

try r = 200 ...
using SVD ...
using NMF ...

```

```

/usr/local/lib/python3.8/dist-packages/sklearn/decomposition/_nmf.py:1637:
ConvergenceWarning: Maximum number of iterations 200 reached. Increase it to
improve convergence.

```

```

warnings.warn(

```

```

using UMAP ...

```

```

[69]: def pkl_read(path, filename):
        with open(path + filename, 'rb') as f:
            # compressed_file = bz2.BZ2File(f, 'r')
            X_dr = pickle.load(f)
        return X_dr
DRs = ['none.pkl',
        'svd_5.pkl', 'svd_20.pkl', 'svd_200.pkl',
        'nmf_5.pkl', 'nmf_20.pkl', 'nmf_200.pkl',
        'umap_5.pkl', 'umap_20.pkl', 'umap_200.pkl']

```

```

[ ]: ks = [10, 20, 50]
km_scores = [ [0] * len(DRs) for i in range(3)]
for i in range(len(ks)):
    print("\nk =", ks[i], "...")
    kmeans_tmp = KMeans(n_clusters=ks[i], init='k-means++', max_iter=2000,
↳ n_init=50, random_state=0)
    for j in range(len(DRs)):
        print("Current dr:", DRs[j])
        # reading Dimensionality Reduction data
        X_dr = pkl_read(path, DRs[j])
        kmeans_fit = kmeans_tmp.fit(X_dr)

```

```
# To increase efficiency, we calculate the average score of each  
→ combination instead.
```

```
tmp_s = 0  
tmp_s += homogeneity_score(y_data, kmeans_fit.labels_)  
tmp_s += completeness_score(y_data, kmeans_fit.labels_)  
tmp_s += v_measure_score(y_data, kmeans_fit.labels_)  
tmp_s += adjusted_rand_score(y_data, kmeans_fit.labels_)  
tmp_s += adjusted_mutual_info_score(y_data, kmeans_fit.labels_)  
tmp_s /= 5  
km_scores[i][j] = tmp_s  
print("Kmeans when k =", ks[i], ", average score:", tmp_s)
```

k = 10 ...

```
Current dr: none.pkl  
Kmeans when k = 10 , average score: 0.3128862783009779  
Current dr: svd_5.pkl  
Kmeans when k = 10 , average score: 0.2882612154581411  
Current dr: svd_20.pkl  
Kmeans when k = 10 , average score: 0.28103632815154544  
Current dr: svd_200.pkl  
Kmeans when k = 10 , average score: 0.2739717062086616  
Current dr: nmf_5.pkl  
Kmeans when k = 10 , average score: 0.24139614802401646  
Current dr: nmf_20.pkl  
Kmeans when k = 10 , average score: 0.21331247059166852  
Current dr: nmf_200.pkl  
Kmeans when k = 10 , average score: 0.032385718417162965  
Current dr: umap_5.pkl  
Kmeans when k = 10 , average score: 0.497940596596295  
Current dr: umap_20.pkl  
Kmeans when k = 10 , average score: 0.5023552400963418  
Current dr: umap_200.pkl  
Kmeans when k = 10 , average score: 0.5034860305394875
```

k = 20 ...

```
Current dr: none.pkl  
Kmeans when k = 20 , average score: 0.319977761965902  
Current dr: svd_5.pkl  
Kmeans when k = 20 , average score: 0.292763515525936  
Current dr: svd_20.pkl  
Kmeans when k = 20 , average score: 0.30654759487815475  
Current dr: svd_200.pkl  
Kmeans when k = 20 , average score: 0.31032718027283984  
Current dr: nmf_5.pkl  
Kmeans when k = 20 , average score: 0.2396382853881612  
Current dr: nmf_20.pkl
```

```

Kmeans when k = 20 , average score: 0.2627086703407877
Current dr: nmf_200.pkl
Kmeans when k = 20 , average score: 0.11487118434220356
Current dr: umap_5.pkl
Kmeans when k = 20 , average score: 0.5641612641269881
Current dr: umap_20.pkl
Kmeans when k = 20 , average score: 0.5488304758703203
Current dr: umap_200.pkl
Kmeans when k = 20 , average score: 0.5568396812973211

```

```

k = 50 ...
Current dr: none.pkl
Kmeans when k = 50 , average score: 0.346192788364497
Current dr: svd_5.pkl
Kmeans when k = 50 , average score: 0.28413792519012615
Current dr: svd_20.pkl
Kmeans when k = 50 , average score: 0.32867074172086935
Current dr: svd_200.pkl
Kmeans when k = 50 , average score: 0.3389048121584812
Current dr: nmf_5.pkl
Kmeans when k = 50 , average score: 0.23368389238229748
Current dr: nmf_20.pkl
Kmeans when k = 50 , average score: 0.3047935570521014
Current dr: nmf_200.pkl
Kmeans when k = 50 , average score: 0.13782058586457885
Current dr: umap_5.pkl
Kmeans when k = 50 , average score: 0.5148044954293607
Current dr: umap_20.pkl
Kmeans when k = 50 , average score: 0.5303504605460957
Current dr: umap_200.pkl
Kmeans when k = 50 , average score: 0.5273297101254659

```

```

[ ]: agg_n = 20
agg_scores = [ [0] * len(DRs)]
for j in range(len(DRs)):
    print("Current dr:", DRs[j])
    # it takes over 1hr to finish the none.pkl, we save the result that we have_
    ↳ found before in advance here
    if (DRs[j] == 'none.pkl'):
        print("Agglomerative Clustering when n_clusters =", agg_n, ", average score:
        ↳ ", 0.3437828455142916)
        agg_scores[0][j] = 0.3437828455142916
        continue
    # reading Dimensionality Reduction data
    X_dr = pickle.read(path, DRs[j])
    # if (DRs[j] == 'none.pkl'):
    #     X_dr = X_dr.toarray()

```

```

agg_tmp = AgglomerativeClustering(n_clusters=agg_n, linkage='ward').fit(X_dr)
# To increase efficiency, we calculate the average score of each combination
→ instead.
tmp_s = 0
tmp_s += homogeneity_score(y_data, agg_tmp.labels_)
tmp_s += completeness_score(y_data, agg_tmp.labels_)
tmp_s += v_measure_score(y_data, agg_tmp.labels_)
tmp_s += adjusted_rand_score(y_data, agg_tmp.labels_)
tmp_s += adjusted_mutual_info_score(y_data, agg_tmp.labels_)
tmp_s /= 5
agg_scores[0][j] = tmp_s
print("Agglomerative Clustering when n_clusters =", agg_n, ", average score:
→", tmp_s)

```

```

Current dr: none.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.3437828455142916
Current dr: svd_5.pkl
Agglomerative Clustering when n_clusters = 20 , average score: 0.281758252396262
Current dr: svd_20.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.35480890471057575
Current dr: svd_200.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.32757379932917774
Current dr: nmf_5.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.23868230288974474
Current dr: nmf_20.pkl
Agglomerative Clustering when n_clusters = 20 , average score: 0.331828295098176
Current dr: nmf_200.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.10838508498003048
Current dr: umap_5.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.5437828303701593
Current dr: umap_20.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.5311649868566655
Current dr: umap_200.pkl
Agglomerative Clustering when n_clusters = 20 , average score:
0.5290589073081226

```

```

[ ]: mcss = [100, 200]
hdb_scores = [ [0] * len(DRs) for i in range(2)]
for i in range(len(mcss)):
    print("\nmin_cluster_size =", mcss[i], "...")

```

```

for j in range(len(DRs)):
    print("Current dr:", DRs[j])
    # reading Dimensionality Reduction data
    X_dr = pickle.load(open(path, DRs[j]))
    y_pred_hdbs = hdbscan.HDBSCAN(min_cluster_size=mcss[i],
    ↪allow_single_cluster=True).fit_predict(X_dr)
    # To increase efficiency, we calculate the average score of each
    ↪combination instead.
    tmp_s = 0
    tmp_s += homogeneity_score(y_data, y_pred_hdbs)
    tmp_s += completeness_score(y_data, y_pred_hdbs)
    tmp_s += v_measure_score(y_data, y_pred_hdbs)
    tmp_s += adjusted_rand_score(y_data, y_pred_hdbs)
    tmp_s += adjusted_mutual_info_score(y_data, y_pred_hdbs)
    tmp_s /= 5
    hdbs_scores[i][j] = tmp_s
    print("HDBSCAN when min_cluster_size =", mcss[i], ", average score:", tmp_s)

```

```

min_cluster_size = 100 ...
Current dr: none.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.016708235408510047
Current dr: svd_5.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.010478759361255884
Current dr: svd_20.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.008450507111460306
Current dr: svd_200.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.006087403063703824
Current dr: nmf_5.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.09425820851516531
Current dr: nmf_20.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.010023000451534418
Current dr: nmf_200.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.006022948696612804
Current dr: umap_5.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.019719014605143435
Current dr: umap_20.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.024785356321964856
Current dr: umap_200.pkl
HDBSCAN when min_cluster_size = 100 , average score: 0.018821756889884196

min_cluster_size = 200 ...
Current dr: none.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.023505928045753802
Current dr: svd_5.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.022856027344854608
Current dr: svd_20.pkl

```

```

HDBSCAN when min_cluster_size = 200 , average score: 0.006763272708372309
Current dr: svd_200.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.007532258670429807
Current dr: nmf_5.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.04213099629994795
Current dr: nmf_20.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.007863958796589979
Current dr: nmf_200.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.005149949351729706
Current dr: umap_5.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.01843515443795917
Current dr: umap_20.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.024252152756405333
Current dr: umap_200.pkl
HDBSCAN when min_cluster_size = 200 , average score: 0.017886948206185305

```

Ans:

From the result we obtained, the combination of **UMAP using ‘cosine’ matrix and n_components setting to 5 plus K-Means using k = 20** works best together for 20-class text data.

It might be because of that the categories it need to predict is also 20. On the other hand, UMAP captures the structure quickly and efficiently comparing to other methods. Combining these two, they perform better than other combinations.

0.19 Question 18

Extra credit: If you can find creative ways to further enhance the clustering performance, report your method and the results you obtain.

Ans:

From the result we obtained, we can see that any clustering method using UMAP as the Dimensionality Reduction method performs much better than other methods.

From this perspective, we try to adjust the parameters of UMAP to see if we can obtain a better result.

```

[ ]: rs18 = [30, 50, 100]
for i in range(len(rs18)):
    print("\ntry r =", rs18[i], "...")

    # UMAP
    print("using UMAP ...")
    umap_tmp = umap.UMAP(n_components=rs18[i], metric='cosine')
    X_umap = umap_tmp.fit_transform(X_tfidf)
    pickle.save(path, X_umap, 'umap_' + str(rs18[i]) + '.pkl')

```

try r = 30 ...

```
using UMAP ...
```

```
try r = 50 ...
```

```
using UMAP ...
```

```
try r = 100 ...
```

```
using UMAP ...
```

```
[ ]: umaps = ['umap_5.pkl', 'umap_20.pkl', 'umap_30.pkl', 'umap_50.pkl', 'umap_100.
↳pkl', 'umap_200.pkl']
```

```
[ ]: ks = [10, 20, 30, 50, 100]
km_scores = [ [0] * len(umaps) for i in range(len(ks))]
for i in range(len(ks)):
    print("\nk =", ks[i], "...")
    kmeans_tmp = KMeans(n_clusters=ks[i], init='k-means++', max_iter=2000,
↳n_init=50, random_state=0)
    for j in range(len(umaps)):
        print("Current UMAP:", umaps[j])
        # reading Dimensionality Reduction data
        X_dr = pkl_read(path, umaps[j])
        kmeans_fit = kmeans_tmp.fit(X_dr)
        # To increase efficiency, we calculate the average score of each
↳combination instead.
        tmp_s = 0
        tmp_s += homogeneity_score(y_data, kmeans_fit.labels_)
        tmp_s += completeness_score(y_data, kmeans_fit.labels_)
        tmp_s += v_measure_score(y_data, kmeans_fit.labels_)
        tmp_s += adjusted_rand_score(y_data, kmeans_fit.labels_)
        tmp_s += adjusted_mutual_info_score(y_data, kmeans_fit.labels_)
        tmp_s /= 5
        km_scores[i][j] = tmp_s
    print("Kmeans when k =", ks[i], ", average score:", tmp_s)
```

```
k = 10 ...
```

```
Current UMAP: umap_5.pkl
```

```
Kmeans when k = 10 , average score: 0.4983759015992576
```

```
Current UMAP: umap_20.pkl
```

```
Kmeans when k = 10 , average score: 0.5023552400963419
```

```
Current UMAP: umap_30.pkl
```

```
Kmeans when k = 10 , average score: 0.5025394557537396
```

```
Current UMAP: umap_50.pkl
```

```
Kmeans when k = 10 , average score: 0.5028938404781135
```

```
Current UMAP: umap_100.pkl
```

```
Kmeans when k = 10 , average score: 0.5012430382777919
```

```
Current UMAP: umap_200.pkl
```

```
Kmeans when k = 10 , average score: 0.5034860305394877
```



```

k = 20 ...
Current UMAP: umap_5.pkl
Kmeans when k = 20 , average score: 0.5527447708674883
Current UMAP: umap_20.pkl
Kmeans when k = 20 , average score: 0.5485179865082734
Current UMAP: umap_30.pkl
Kmeans when k = 20 , average score: 0.5594055845292851
Current UMAP: umap_50.pkl
Kmeans when k = 20 , average score: 0.5648750854798654
Current UMAP: umap_100.pkl
Kmeans when k = 20 , average score: 0.5660340490887646
Current UMAP: umap_200.pkl
Kmeans when k = 20 , average score: 0.556148359746299

k = 30 ...
Current UMAP: umap_5.pkl
Kmeans when k = 30 , average score: 0.5570951177263229
Current UMAP: umap_20.pkl
Kmeans when k = 30 , average score: 0.5551454175119898
Current UMAP: umap_30.pkl
Kmeans when k = 30 , average score: 0.5536573120116477
Current UMAP: umap_50.pkl
Kmeans when k = 30 , average score: 0.5495626680424508
Current UMAP: umap_100.pkl
Kmeans when k = 30 , average score: 0.5503191097974127
Current UMAP: umap_200.pkl
Kmeans when k = 30 , average score: 0.559685131958972

k = 50 ...
Current UMAP: umap_5.pkl
Kmeans when k = 50 , average score: 0.5225005310003482
Current UMAP: umap_20.pkl
Kmeans when k = 50 , average score: 0.5311132275419802
Current UMAP: umap_30.pkl
Kmeans when k = 50 , average score: 0.5165038124127054
Current UMAP: umap_50.pkl
Kmeans when k = 50 , average score: 0.520842165533297
Current UMAP: umap_100.pkl
Kmeans when k = 50 , average score: 0.51853794581093
Current UMAP: umap_200.pkl
Kmeans when k = 50 , average score: 0.5255663945896251

k = 100 ...
Current UMAP: umap_5.pkl
Kmeans when k = 100 , average score: 0.4731704103593664
Current UMAP: umap_20.pkl
Kmeans when k = 100 , average score: 0.480518314730074

```

```

Current UMAP: umap_30.pkl
Kmeans when k = 100 , average score: 0.4752450376025603
Current UMAP: umap_50.pkl
Kmeans when k = 100 , average score: 0.48481560206411894
Current UMAP: umap_100.pkl
Kmeans when k = 100 , average score: 0.47507179317347054
Current UMAP: umap_200.pkl
Kmeans when k = 100 , average score: 0.478029997188872

```

From above result, we find that the performace of UMAP_100 is a little bit better than UMAP_50. And $k = 20$ is the best setting for K-Means. We will try to adjust some parameters for these two settings to get a better result.

```

[ ]: n_neighs = [20, 30, 50, 100]
min_dists = [0.2, 0.3, 0.5, 0.87]
kmeans_tmp = KMeans(n_clusters=20, init='k-means++', max_iter=2000, n_init=50,
    random_state=0)
for i in range(len(n_neighs)):
    print("\nn_neighbors =", n_neighs[i], "...")
    for j in range(len(min_dists)):
        print("Current min_dist:", min_dists[j])
        umap_tmp = umap.UMAP(n_components=100, n_neighbors=n_neighs[i],
    min_dist=min_dists[j], metric='cosine')
        X_umap = umap_tmp.fit_transform(X_tfidf)
        kmeans_fit = kmeans_tmp.fit(X_umap)
        tmp_s = 0
        tmp_s += homogeneity_score(y_data, kmeans_fit.labels_)
        tmp_s += completeness_score(y_data, kmeans_fit.labels_)
        tmp_s += v_measure_score(y_data, kmeans_fit.labels_)
        tmp_s += adjusted_rand_score(y_data, kmeans_fit.labels_)
        tmp_s += adjusted_mutual_info_score(y_data, kmeans_fit.labels_)
        tmp_s /= 5
        km_scores[i][j] = tmp_s
        print("Kmeans when k = 20", ", UMAP when n_neighbors =", n_neighs[i],
    min_dist = min_dists[j], ", average score:", tmp_s)

```

```

n_neighbors = 20 ...
Current min_dist: 0.2
Kmeans when k = 20 , UMAP when n_neighbors = 20 , min_dist = 0.2 , average
score: 0.580787761743627
Current min_dist: 0.3
Kmeans when k = 20 , UMAP when n_neighbors = 20 , min_dist = 0.3 , average
score: 0.5724508070219805
Current min_dist: 0.5
Kmeans when k = 20 , UMAP when n_neighbors = 20 , min_dist = 0.5 , average
score: 0.5744771322175886
Current min_dist: 0.87

```

Kmeans when k = 20 , UMAP when n_neighbors = 20 , min_dist = 0.87 , average score: 0.5945454230665579

n_neighbors = 30 ...

Current min_dist: 0.2

Kmeans when k = 20 , UMAP when n_neighbors = 30 , min_dist = 0.2 , average score: 0.5737626120600483

Current min_dist: 0.3

Kmeans when k = 20 , UMAP when n_neighbors = 30 , min_dist = 0.3 , average score: 0.5736243636591776

Current min_dist: 0.5

Kmeans when k = 20 , UMAP when n_neighbors = 30 , min_dist = 0.5 , average score: 0.5778039370853855

Current min_dist: 0.87

Kmeans when k = 20 , UMAP when n_neighbors = 30 , min_dist = 0.87 , average score: 0.595183451482399

n_neighbors = 50 ...

Current min_dist: 0.2

Kmeans when k = 20 , UMAP when n_neighbors = 50 , min_dist = 0.2 , average score: 0.576943355602227

Current min_dist: 0.3

Kmeans when k = 20 , UMAP when n_neighbors = 50 , min_dist = 0.3 , average score: 0.5755555116333844

Current min_dist: 0.5

Kmeans when k = 20 , UMAP when n_neighbors = 50 , min_dist = 0.5 , average score: 0.5911223311681624

Current min_dist: 0.87

Kmeans when k = 20 , UMAP when n_neighbors = 50 , min_dist = 0.87 , average score: 0.5970547503057227

n_neighbors = 100 ...

Current min_dist: 0.2

Kmeans when k = 20 , UMAP when n_neighbors = 100 , min_dist = 0.2 , average score: 0.5825166086844924

Current min_dist: 0.3

Kmeans when k = 20 , UMAP when n_neighbors = 100 , min_dist = 0.3 , average score: 0.5785838926897937

Current min_dist: 0.5

Kmeans when k = 20 , UMAP when n_neighbors = 100 , min_dist = 0.5 , average score: 0.5798158369565941

Current min_dist: 0.87

Kmeans when k = 20 , UMAP when n_neighbors = 100 , min_dist = 0.87 , average score: 0.5967891400756058

After the experiment, we found that the performance usually gets better when we increase the **min_dist** parameter. Which makes sense because as the minimum distance between each data point increase, we should separate the cluster more easily and get a higher Homogeneity and

Completeness score.

In conclusion, we found a better clustering performance by using **UMAP with \$ n_components = 100, n_neighbors = 50\$, and $min_dist = 0.87$. Combining it with K-means with $k = 20$** , we get a average score which is approximately 0.597.

0.20 Question 19

In a brief paragraph discuss: If the VGG network is trained on a dataset with perhaps totally different classes as targets, why would one expect the features derived from such a network to have discriminative power for a custom dataset?

Ans:

We expect that the model is trained on a large and general enough dataset. Even though it's a dataset with perhaps totally different classes as targets, we can still get advantage from its learned feature maps to repurpose it.

To be specifically, there are always multiple layers inside our network model. We can imagine that the final layer of our network is the one that learn to identify classes specific to our project that need training. The initial layers or the middle layers are used to detect slant lines no matter what you want to classify. Therefore, we can always reuse the neural network instead of retraining them every time we create a new one.

Finally, after getting the pretrained model, we can use fine-tuning to make the model more relevant to the custom dataset we want to use. By doing this, we can expect that the features derived from previous network will have discriminative power for a custom dataset.

0.21 Question 20

In a brief paragraph explain how the helper code base is performing feature extraction.

Ans:

It first extracts feature layers and pooling layer from the VGG-16. After that, it constructs a flatten layer using torch, and extracts the first part of fully-connected layer from VGG-16.

Similar to the concept we mentioned in Q19, it extracts the basic layers from the pretrained model VGG-16. By doing this, it extracts the basic discriminative power we have in the pretrained model. We can use it in further application such as fine-tuning on the custom dataset we want to test.

0.22 Question 21

```
[4]: filename = './flowers_features_and_labels.npz'

if os.path.exists(filename):
    file = np.load(filename)
    f_all, y_all = file['f_all'], file['y_all']

else:
    if not os.path.exists('./flower_photos'):
        # download the flowers dataset and extract its images
```

```

url = 'http://download.tensorflow.org/example_images/flower_photos.tgz'
with open('./flower_photos.tgz', 'wb') as file:
    file.write(requests.get(url).content)
with tarfile.open('./flower_photos.tgz') as file:
    file.extractall('.')
os.remove('./flower_photos.tgz')

class FeatureExtractor(nn.Module):
    def __init__(self):
        super().__init__()

        vgg = torch.hub.load('pytorch/vision:v0.10.0', 'vgg16',
→pretrained=True)

        # Extract VGG-16 Feature Layers
        self.features = list(vgg.features)
        self.features = nn.Sequential(*self.features)
        # Extract VGG-16 Average Pooling Layer
        self.pooling = vgg.avgpool
        # Convert the image into one-dimensional vector
        self.flatten = nn.Flatten()
        # Extract the first part of fully-connected layer from VGG16
        self.fc = vgg.classifier[0]

    def forward(self, x):
        # It will take the input 'x' until it returns the feature vector
→called 'out'
        out = self.features(x)
        out = self.pooling(out)
        out = self.flatten(out)
        out = self.fc(out)
        return out

# Initialize the model
assert torch.cuda.is_available()
feature_extractor = FeatureExtractor().cuda().eval()

dataset = datasets.ImageFolder(root='./flower_photos',
                                transform=transforms.Compose([transforms.
→Resize(224),
                                                                transforms.
→CenterCrop(224),
                                                                transforms.
→ToTensor(),
                                                                transforms.
→Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])]))

```

```
dataloader = DataLoader(dataset, batch_size=64, shuffle=True)

# Extract features and store them on disk
f_all, y_all = np.zeros((0, 4096)), np.zeros((0,))
for x, y in tqdm(dataloader):
    with torch.no_grad():
        f_all = np.vstack([f_all, feature_extractor(x.cuda()).cpu()])
        y_all = np.concatenate([y_all, y])
np.savez(filename, f_all=f_all, y_all=y_all)
```

Downloading: "https://github.com/pytorch/vision/zipball/v0.10.0" to
 /root/.cache/torch/hub/v0.10.0.zip
 /usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:208:
 UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
 removed in the future, please use 'weights' instead.
 warnings.warn(
 /usr/local/lib/python3.8/dist-packages/torchvision/models/_utils.py:223:
 UserWarning: Arguments other than a weight enum or `None` for 'weights' are
 deprecated since 0.13 and may be removed in the future. The current behavior is
 equivalent to passing `weights=VGG16_Weights.IMAGENET1K_V1`. You can also use
 `weights=VGG16_Weights.DEFAULT` to get the most up-to-date weights.
 warnings.warn(msg)
 Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to
 /root/.cache/torch/hub/checkpoints/vgg16-397923af.pth

```
0%|          | 0.00/528M [00:00<?, ?B/s]
```

```
100%|         | 58/58 [00:44<00:00, 1.32it/s]
```

How many pixels are there in the original images? How many features does the VGG network extract per image; i.e what is the dimension of each feature vector for an image sample?

Ans:

There are 224 x 224 pixels in the original images.

There are totally 3670 features being extracted by the VGG network per image.

```
[ ]: for x, y in tqdm(dataloader):
    print("\n")
    print(x.size())
    print(y.size())
    break

print("\nfeature vector:", f_all.shape, y_all.shape)
num_features = f_all.shape[1]
```

```
0%|          | 0/58 [00:00<?, ?it/s]
```

```
torch.Size([64, 3, 224, 224])
torch.Size([64])
```

feature vector: (3670, 4096) (3670,)

0.23 Question 22

Are the extracted features dense or sparse? (Compare with sparse TF-IDF features in text.)

Ans:

From Q10, we found the shape of TFIDF dataset: (18846, 45365)

The extracted feature vectors we found: (3670, 4096)

Furthermore, we can see that the numbers inside TF-IDF features are much closer to zero comparing to the extracted features. Therefore, the extracted features are more dense.

```
[ ]: print(f_all.shape)
      print(f_all)
      print(X_tfidf.shape)
      print(X_tfidf)

(3670, 4096)
[[-0.45783606  2.75369859  0.84992045 ... -2.20686865 -0.35974157
   1.0510205 ]
 [-0.01870249 -2.30375719 -0.63398772 ... -4.68432951  1.58356237
   2.4888947 ]
 [-2.29198527 -1.39408898  0.04004309 ... -1.57638204 -0.39932269
  -0.77769351]
 ...
 [-1.71726406 -3.87869978 -4.15782833 ... -0.03698874  0.24791169
   7.03038645]
 [-1.67345393 -0.80352944 -2.1430831  ...  2.40008855 -0.81528664
   1.85219288]
 [ 1.27578044 -0.57890528 -0.29694718 ... -3.28786254 -1.68270898
   0.94054085]]
(18846, 45365)
(0, 45082)    0.10834863946498498
(0, 44381)    0.1584590937040802
(0, 42541)    0.058066446887881366
(0, 41490)    0.08051514880421468
(0, 41316)    0.13329293384982419
(0, 38204)    0.16925886266575615
(0, 37492)    0.09850240968265686
(0, 36007)    0.09826054436567758
(0, 33328)    0.47942224155915814
(0, 33262)    0.18420734272692216
```

| | |
|----------------|----------------------|
| (0, 31695) | 0.3963165278889442 |
| (0, 31683) | 0.19522385821708108 |
| (0, 29617) | 0.11875524902913023 |
| (0, 29022) | 0.0622296687474105 |
| (0, 27706) | 0.11540121691934306 |
| (0, 25516) | 0.08157968927565735 |
| (0, 24419) | 0.11312900170774419 |
| (0, 24371) | 0.1181419537275807 |
| (0, 22243) | 0.2851039251453595 |
| (0, 18123) | 0.16532689809474904 |
| (0, 17976) | 0.1311636085943422 |
| (0, 17787) | 0.1240128023295206 |
| (0, 15208) | 0.13079807275779112 |
| (0, 14150) | 0.20396815849564107 |
| (0, 11697) | 0.09878805754148298 |
| : | : |
| (18845, 18514) | 0.05860474143465704 |
| (18845, 17963) | 0.05712310612173224 |
| (18845, 17506) | 0.0722870974794078 |
| (18845, 17232) | 0.05272356612438965 |
| (18845, 16972) | 0.06351235739684634 |
| (18845, 16255) | 0.09477122354699491 |
| (18845, 14717) | 0.1514142823050571 |
| (18845, 14608) | 0.16023615155667814 |
| (18845, 13592) | 0.04186755908990516 |
| (18845, 13576) | 0.08863063242402147 |
| (18845, 13168) | 0.048522897301106066 |
| (18845, 12635) | 0.060843936981812494 |
| (18845, 12514) | 0.13454853488766122 |
| (18845, 11547) | 0.20054431862465735 |
| (18845, 9782) | 0.08533659547734827 |
| (18845, 9006) | 0.15577532500269423 |
| (18845, 8379) | 0.0701110751954832 |
| (18845, 7355) | 0.08376382234918138 |
| (18845, 7162) | 0.07120001286865578 |
| (18845, 6709) | 0.06481462637628999 |
| (18845, 5388) | 0.06852361213473072 |
| (18845, 5385) | 0.050512997363727795 |
| (18845, 5154) | 0.09822577469614639 |
| (18845, 2468) | 0.0792566940653086 |
| (18845, 826) | 0.04868629510260459 |

0.24 Question 23

In order to inspect the high-dimensional features, t-SNE is a popular off-the-shelf choice for visualizing Vision features. Map the features you have extracted onto 2 dimensions with t-SNE. Then plot the mapped feature vectors along x and y axes. Color-code the data points with ground-truth labels. Describe your observation.

Ans:

From our result, we can see that there are 5 categories (clusterings) and they have many overlapping spots when we project it onto 2 dimensions.

It indicates that we might face some problems when we try to use high dimensional clustering method on this dataset. The distance between each data point is really close according to our projection result.

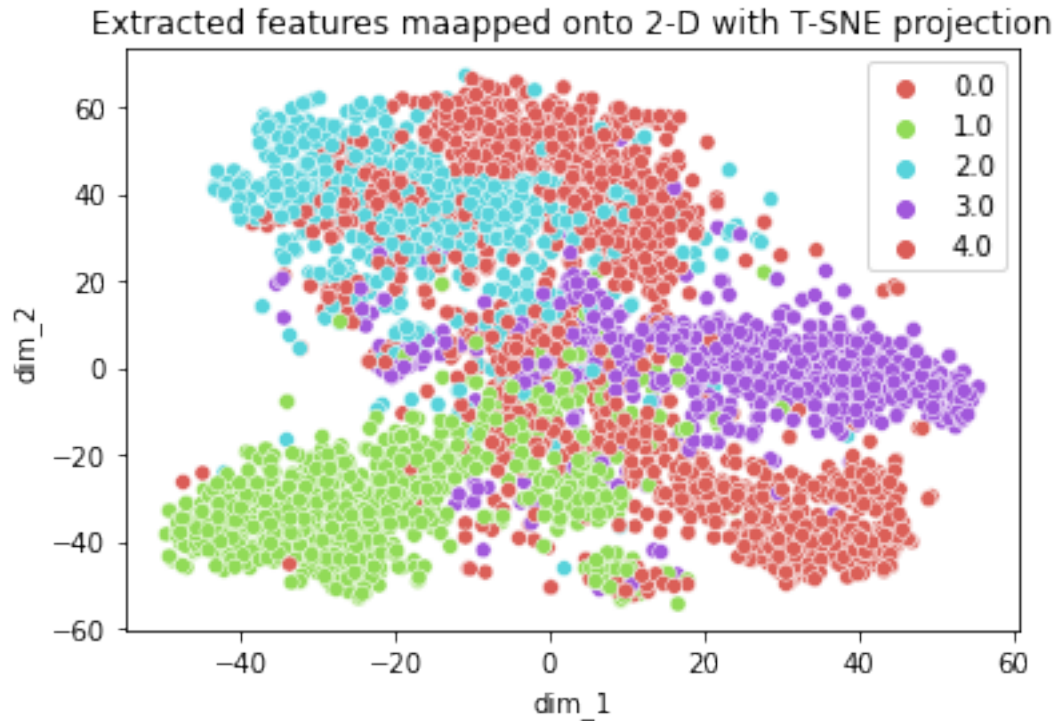
```
[ ]: tsne = TSNE(n_components=2, verbose=1, random_state=0)
      z = tsne.fit_transform(f_all)
      df = pd.DataFrame()
      df["y"] = y_all
      df["dim_1"] = z[:,0]
      df["dim_2"] = z[:,1]
      sns.scatterplot(x="dim_1", y="dim_2", hue=df.y.tolist(),
                     palette=sns.color_palette("hls", as_cmap = True),
                     data=df).set(title="Extracted features maapped onto 2-D with
↳T-SNE projection")
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/manifold/_t_sne.py:780:
FutureWarning: The default initialization in TSNE will change from 'random' to
'pca' in 1.2.
```

```
warnings.warn(
/usr/local/lib/python3.8/dist-packages/sklearn/manifold/_t_sne.py:790:
FutureWarning: The default learning rate in TSNE will change from 200.0 to
'auto' in 1.2.
warnings.warn(
```

```
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 3670 samples in 0.018s...
[t-SNE] Computed neighbors for 3670 samples in 3.112s...
[t-SNE] Computed conditional probabilities for sample 1000 / 3670
[t-SNE] Computed conditional probabilities for sample 2000 / 3670
[t-SNE] Computed conditional probabilities for sample 3000 / 3670
[t-SNE] Computed conditional probabilities for sample 3670 / 3670
[t-SNE] Mean sigma: 36.158544
[t-SNE] KL divergence after 250 iterations with early exaggeration: 78.663742
[t-SNE] KL divergence after 1000 iterations: 1.762485
```

```
[ ]: [Text(0.5, 1.0, 'Extracted features maapped onto 2-D with T-SNE projection')]
```



0.25 Question 24

0.25.1 Autoencoder

```
[5]: class Autoencoder(torch.nn.Module, TransformerMixin):
    def __init__(self, n_components):
        super().__init__()
        self.n_components = n_components
        self.n_features = None # to be determined with data
        self.encoder = None
        self.decoder = None

    def _create_encoder(self):
        return nn.Sequential(
            nn.Linear(4096, 1280),
            nn.ReLU(True),
            nn.Linear(1280, 640),
            nn.ReLU(True), nn.Linear(640, 120), nn.ReLU(True), nn.Linear(120,
↪self.n_components))

    def _create_decoder(self):
        return nn.Sequential(
            nn.Linear(self.n_components, 120),
            nn.ReLU(True),
```

```

        nn.Linear(120, 640),
        nn.ReLU(True),
        nn.Linear(640, 1280),
        nn.ReLU(True), nn.Linear(1280, 4096))

def forward(self, X):
    encoded = self.encoder(X)
    decoded = self.decoder(encoded)
    return decoded

def fit(self, X):
    X = torch.tensor(X, dtype=torch.float32, device='cuda')
    self.n_features = X.shape[1]
    self.encoder = self._create_encoder()
    self.decoder = self._create_decoder()
    self.cuda()
    self.train()

    criterion = nn.MSELoss()
    optimizer = torch.optim.Adam(self.parameters(), lr=1e-3,
    ↪weight_decay=1e-5)

    dataset = TensorDataset(X)
    dataloader = DataLoader(dataset, batch_size=128, shuffle=True)

    for epoch in tqdm(range(100)):
        for (X_,) in dataloader:
            X_ = X_.cuda()
            # =====forward=====
            output = self(X_)
            loss = criterion(output, X_)
            # =====backward=====
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    return self

def transform(self, X):
    X = torch.tensor(X, dtype=torch.float32, device='cuda')
    self.eval()
    with torch.no_grad():
        return self.encoder(X).cpu().numpy()

```

Report the best result (in terms of rand score) within the table below. For HDBSCAN, introduce a conservative parameter grid over min cluster size and min samples.

| Module | Alternatives | Hyperparameters |
|--------------------------|--------------------------|---------------------------------------|
| Dimensionality Reduction | None | N/A |
| Dimensionality Reduction | SVD | $r = 50$ |
| Dimensionality Reduction | UMAP | $n_components = 50$ |
| Dimensionality Reduction | Autoencoder | $num_features = 50$ |
| Clustering | K-Means | $k = 5$ |
| Clustering | Agglomerative Clustering | $n_clusters = 5$ |
| Clustering | HDBSCAN | $min_cluster_size$ & $min_samples$ |

Ans:

The best rand score we get is approximately 0.467. Which is derived from the combination of UMAP [$n_components = 50$] + K-Means [$k = 5$].

For HDBSCAN, we set $min_cluster_size = [50, 100, 200]$, $min_samples = [20, 50, 100, 200]$ to proceed a grid search. However, the best rand score we can get is approximately only 0.094 when combining with UMAP [$n_components = 50$]. Which is still much worse than other combination.

```
[ ]: # None
pkl_save(path, f_all, 'f_none.pkl')

r24 = 50
print("\ntry r =", r24, "...")
# SVD
print("using SVD ...")
svd_tmp = TruncatedSVD(n_components=r24, random_state=0)
X_svd = svd_tmp.fit_transform(f_all)
pkl_save(path, X_svd, 'f_svd_' + str(r24) + '.pkl')

# UMAP
print("using UMAP ...")
umap_tmp = umap.UMAP(n_components=r24, metric='cosine')
X_umap = umap_tmp.fit_transform(f_all)
pkl_save(path, X_umap, 'f_umap_' + str(r24) + '.pkl')

# Autoencoder
print("using Autoencoder ...")
X_aec = Autoencoder(r24).fit_transform(f_all)
pkl_save(path, X_aec, 'f_aec_' + str(r24) + '.pkl')
```

```
try r = 50 ...
using SVD ...
using UMAP ...
using Autoencoder ...

100%|          | 100/100 [00:24<00:00, 4.08it/s]
```

```
[ ]: DRs = ['f_none.pkl', 'f_svd_50.pkl', 'f_umap_50.pkl', 'f_aec_50.pkl']
```

```
[ ]: ks = [5]
km_scores24 = [ [0] * len(DRs) for i in range(len(ks))]
for i in range(len(ks)):
    print("\nk =", ks[i], "...")
    kmeans_tmp = KMeans(n_clusters=ks[i], init='k-means++', max_iter=2000,
    ↪n_init=50, random_state=0)
    for j in range(len(DRs)):
        print("Current dr:", DRs[j])
        # reading Dimensionality Reduction data
        X_dr = pkl_read(path, DRs[j])
        kmeans_fit = kmeans_tmp.fit(X_dr)
        # As the question mentioned, we use rand score to determine which
    ↪combination performs best.
        tmp_s = adjusted_rand_score(y_all, kmeans_fit.labels_)
        km_scores24[i][j] = tmp_s
        print("Kmeans when k =", ks[i], ", rand score:", tmp_s)
```

```
k = 5 ...
Current dr: f_none.pkl
Kmeans when k = 5 , rand score: 0.19216442375229253
Current dr: f_svd_50.pkl
Kmeans when k = 5 , rand score: 0.18805072438502163
Current dr: f_umap_50.pkl
Kmeans when k = 5 , rand score: 0.4665092104522564
Current dr: f_aec_50.pkl
Kmeans when k = 5 , rand score: 0.20346920248488434
```

```
[ ]: agg_n = 5
agg_scores24 = [ [0] * len(DRs)]
for j in range(len(DRs)):
    print("Current dr:", DRs[j])
    # reading Dimensionality Reduction data
    X_dr = pkl_read(path, DRs[j])
    agg_tmp = AgglomerativeClustering(n_clusters=agg_n, linkage='ward').fit(X_dr)
    # As the question mentioned, we use rand score to determine which combination
    ↪performs best.
    tmp_s = adjusted_rand_score(y_all, agg_tmp.labels_)
    agg_scores24[0][j] = tmp_s
    print("Agglomerative Clustering when n_clusters =", agg_n, ", rand score:",
    ↪tmp_s)
```

```
Current dr: f_none.pkl
Agglomerative Clustering when n_clusters = 5 , rand score: 0.18855278251971858
Current dr: f_svd_50.pkl
```

Agglomerative Clustering when n_clusters = 5 , rand score: 0.19998277233193223
 Current dr: f_umap_50.pkl
 Agglomerative Clustering when n_clusters = 5 , rand score: 0.45265411323193355
 Current dr: f_aec_50.pkl
 Agglomerative Clustering when n_clusters = 5 , rand score: 0.28906349861270764

```
[ ]: mc_ss = [50, 100, 200]
min_ss = [20, 50, 100, 200]
# hdb_scores24 = [ [0] * len(DRs) for i in range(2)]
for i in range(len(mc_ss)):
    print("\nmin_cluster_size =", mc_ss[i], "...")
    for k in range(len(min_ss)):
        print("\nmin_samples =", min_ss[k], "...")
        for j in range(len(DRs)):
            print("Current dr:", DRs[j])
            # reading Dimensionality Reduction data
            X_dr = pkl_read(path, DRs[j])
            y_pred_hdb = hdbscan.HDBSCAN(min_cluster_size=mc_ss[i],
            ↪min_samples=min_ss[k], allow_single_cluster=True).fit_predict(X_dr)
            # As the question mentioned, we use rand score to determine which
            ↪combination performs best.
            tmp_s = adjusted_rand_score(y_all, y_pred_hdb)
            # hdb_scores24[i][j] = tmp_s
            print("HDBSCAN when min_cluster_size =", mc_ss[i], ", min_samples =",
            ↪min_ss[k], ", rand score:", tmp_s)
```

min_cluster_size = 50 ...

min_samples = 20 ...

Current dr: f_none.pkl

HDBSCAN when min_cluster_size = 50 , min_samples = 20 , rand score:
 -0.0010448965393208944

Current dr: f_svd_50.pkl

HDBSCAN when min_cluster_size = 50 , min_samples = 20 , rand score:
 -0.000779608175723694

Current dr: f_umap_50.pkl

HDBSCAN when min_cluster_size = 50 , min_samples = 20 , rand score:
 0.09411014415380671

Current dr: f_aec_50.pkl

HDBSCAN when min_cluster_size = 50 , min_samples = 20 , rand score:
 -0.0004412418123981401

min_samples = 50 ...

Current dr: f_none.pkl

HDBSCAN when min_cluster_size = 50 , min_samples = 50 , rand score:
 -0.0007876010819439827

Current dr: f_svd_50.pkl

```

HDBSCAN when min_cluster_size = 50 , min_samples = 50 , rand score:
-0.0006014805513858321
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 50 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 50 , rand score:
-0.0004439061144715697

min_samples = 100 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 100 , rand score:
-0.0009604001021349855
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 100 , rand score:
-0.0006916862073005186
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 100 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 100 , rand score:
-0.0005748375306515365

min_samples = 200 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 200 , rand score:
-0.0010403291643378723
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 200 , rand score:
-0.0006349746345946608
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 200 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 50 , min_samples = 200 , rand score:
-0.0007579131445543391

min_cluster_size = 100 ...

min_samples = 20 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 20 , rand score:
-0.0014971865521274726
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 20 , rand score:
-0.0011620618099798916
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 20 , rand score:

```

```

0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 20 , rand score:
-0.0011309957852153418

min_samples = 50 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 50 , rand score:
-0.0012971990177056832
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 50 , rand score:
-0.0010401276627790335
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 50 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 50 , rand score:
-0.0010012951318233462

min_samples = 100 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 100 , rand score:
-0.0011500237253836286
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 100 , rand score:
-0.0010273129275636568
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 100 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 100 , rand score:
-0.0007411171744202416

min_samples = 200 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 200 , rand score:
-0.0011550719544078678
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 200 , rand score:
-0.0011663333883850172
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 200 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 100 , min_samples = 200 , rand score:
-0.0007966476936868744

min_cluster_size = 200 ...

```



```

min_samples = 20 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 20 , rand score:
-0.0017188731500706628
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 20 , rand score:
-0.0019216794456046568
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 20 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 20 , rand score:
-0.001526570766257553

min_samples = 50 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 50 , rand score:
-0.001625953931001303
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 50 , rand score:
-0.001632013880071044
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 50 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 50 , rand score:
-0.0015944421958386506

min_samples = 100 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 100 , rand score:
-0.0016926133707684525
Current dr: f_svd_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 100 , rand score:
-0.0015988861584897939
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 100 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 100 , rand score:
-0.0014041597950487877

min_samples = 200 ...
Current dr: f_none.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 200 , rand score:
-0.0015665664301178425
Current dr: f_svd_50.pkl

```

```

HDBSCAN when min_cluster_size = 200 , min_samples = 200 , rand score:
-0.0017047332689079343
Current dr: f_umap_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 200 , rand score:
0.09411014415380671
Current dr: f_aec_50.pkl
HDBSCAN when min_cluster_size = 200 , min_samples = 200 , rand score:
-0.0014094117509092297

```

0.26 Question 25

0.26.1 MLP classifier

```

[61]: class MLP(torch.nn.Module):
    def __init__(self, num_features):
        super().__init__()
        self.model = nn.Sequential(
            nn.Linear(num_features, 1280),
            nn.ReLU(True),
            nn.Linear(1280, 640),
            nn.ReLU(True),
            nn.Linear(640, 5),
            nn.LogSoftmax(dim=1)
        )
        self.cuda()

    def forward(self, X):
        return self.model(X)

    def train(self, X, y):
        X = torch.tensor(X, dtype=torch.float32, device='cuda')
        y = torch.tensor(y, dtype=torch.int64, device='cuda')

        self.model.train()

        criterion = nn.NLLLoss()
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3,
↪weight_decay=1e-5)

        dataset = TensorDataset(X, y)
        dataloader = DataLoader(dataset, batch_size=128, shuffle=True)

        for epoch in tqdm(range(100)):
            for (X_, y_) in dataloader:
                X_ = X_.cuda()
                y_ = y_.cuda()
                # =====forward=====

```

```

        output = self(X_)
        loss = criterion(output, y_)
        # =====backward=====
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    return self

def eval(self, X_test, y_test):
    X = torch.tensor(X_test, dtype=torch.float32, device='cuda')
    y = torch.tensor(y_test, dtype=torch.int64, device='cuda')
    dataset = TensorDataset(X, y)
    dataloader = DataLoader(dataset, batch_size=128, shuffle = False)
    correct_pred, num_examples = 0, 0
    for (X_, y_) in dataloader:
        X_ = X_.cuda()
        y_ = y_.cuda()
        logits = self(X_)
        predicted_labels = torch.argmax(logits, 1)
        num_examples += y_.size(0)
        correct_pred += (predicted_labels == y_).sum()
        # print("predicted_labels", predicted_labels)
        # print("y", y_)
    acc = correct_pred.float() / num_examples * 100
    return acc

```

Report the test accuracy of the MLP classifier on the original VGG features.

```

[62]: X_train, X_test, y_train, y_test = train_test_split(f_all, y_all, test_size=0.
      ↪2, random_state=0)
      print(X_train.shape)

```

(2936, 4096)

```

[65]: model = MLP(X_train.shape[1])
      model.train(X_train, y_train)
      acc = model.eval(X_test, y_test)
      print("\nMLP classifier on the original VGG features")
      print("Test accuracy: %.2f" % acc, "%")

```

100%| | 100/100 [00:08<00:00, 11.75it/s]

MLP classifier on the original VGG features
Test accuracy: 90.46 %

Report the same when using the reduced-dimension features (you have freedom in choosing the dimensionality reduction algorithm and its parameters).

```
[73]: umap_50 = umap.UMAP(n_components=50, metric='cosine')
X_umap = umap_50.fit_transform(f_all)
# X_umap = pickle_read(path, 'f_umap_50.pkl')
X_train, X_test, y_train, y_test = train_test_split(X_umap, y_all, test_size=0.
↳2, random_state=0)
print(X_train.shape)
```

(2936, 50)

```
[74]: model = MLP(X_train.shape[1])
model.train(X_train, y_train)
acc = model.eval(X_test, y_test)
print("\nMLP classifier on the reduced-dimension features")
print("Test accuracy: %.2f" % acc, "%")
```

100%| | 100/100 [00:12<00:00, 8.21it/s]

MLP classifier on the reduced-dimension features

Test accuracy: 86.38 %

Does the performance of the model suffer with the reduced-dimension representations? Is it significant?

Ans:

Yes, the performance of the model gets worse when we use the reduced-dimension representation. It is not significant, it only decreases the test accuracy by 4% approximately.

Does the success in classification make sense in the context of the clustering results obtained for the same features in Question 24.

Ans:

Yes, it makes sense. The performance of the model should not be affected a lot if it's already a well-clustering dataset. Even though we reduce the dimensions of the features, the model can still get the structure of the data and performs well eventually.